

Reinforcement Learning

Lab #6: Convolutional Neural Networks

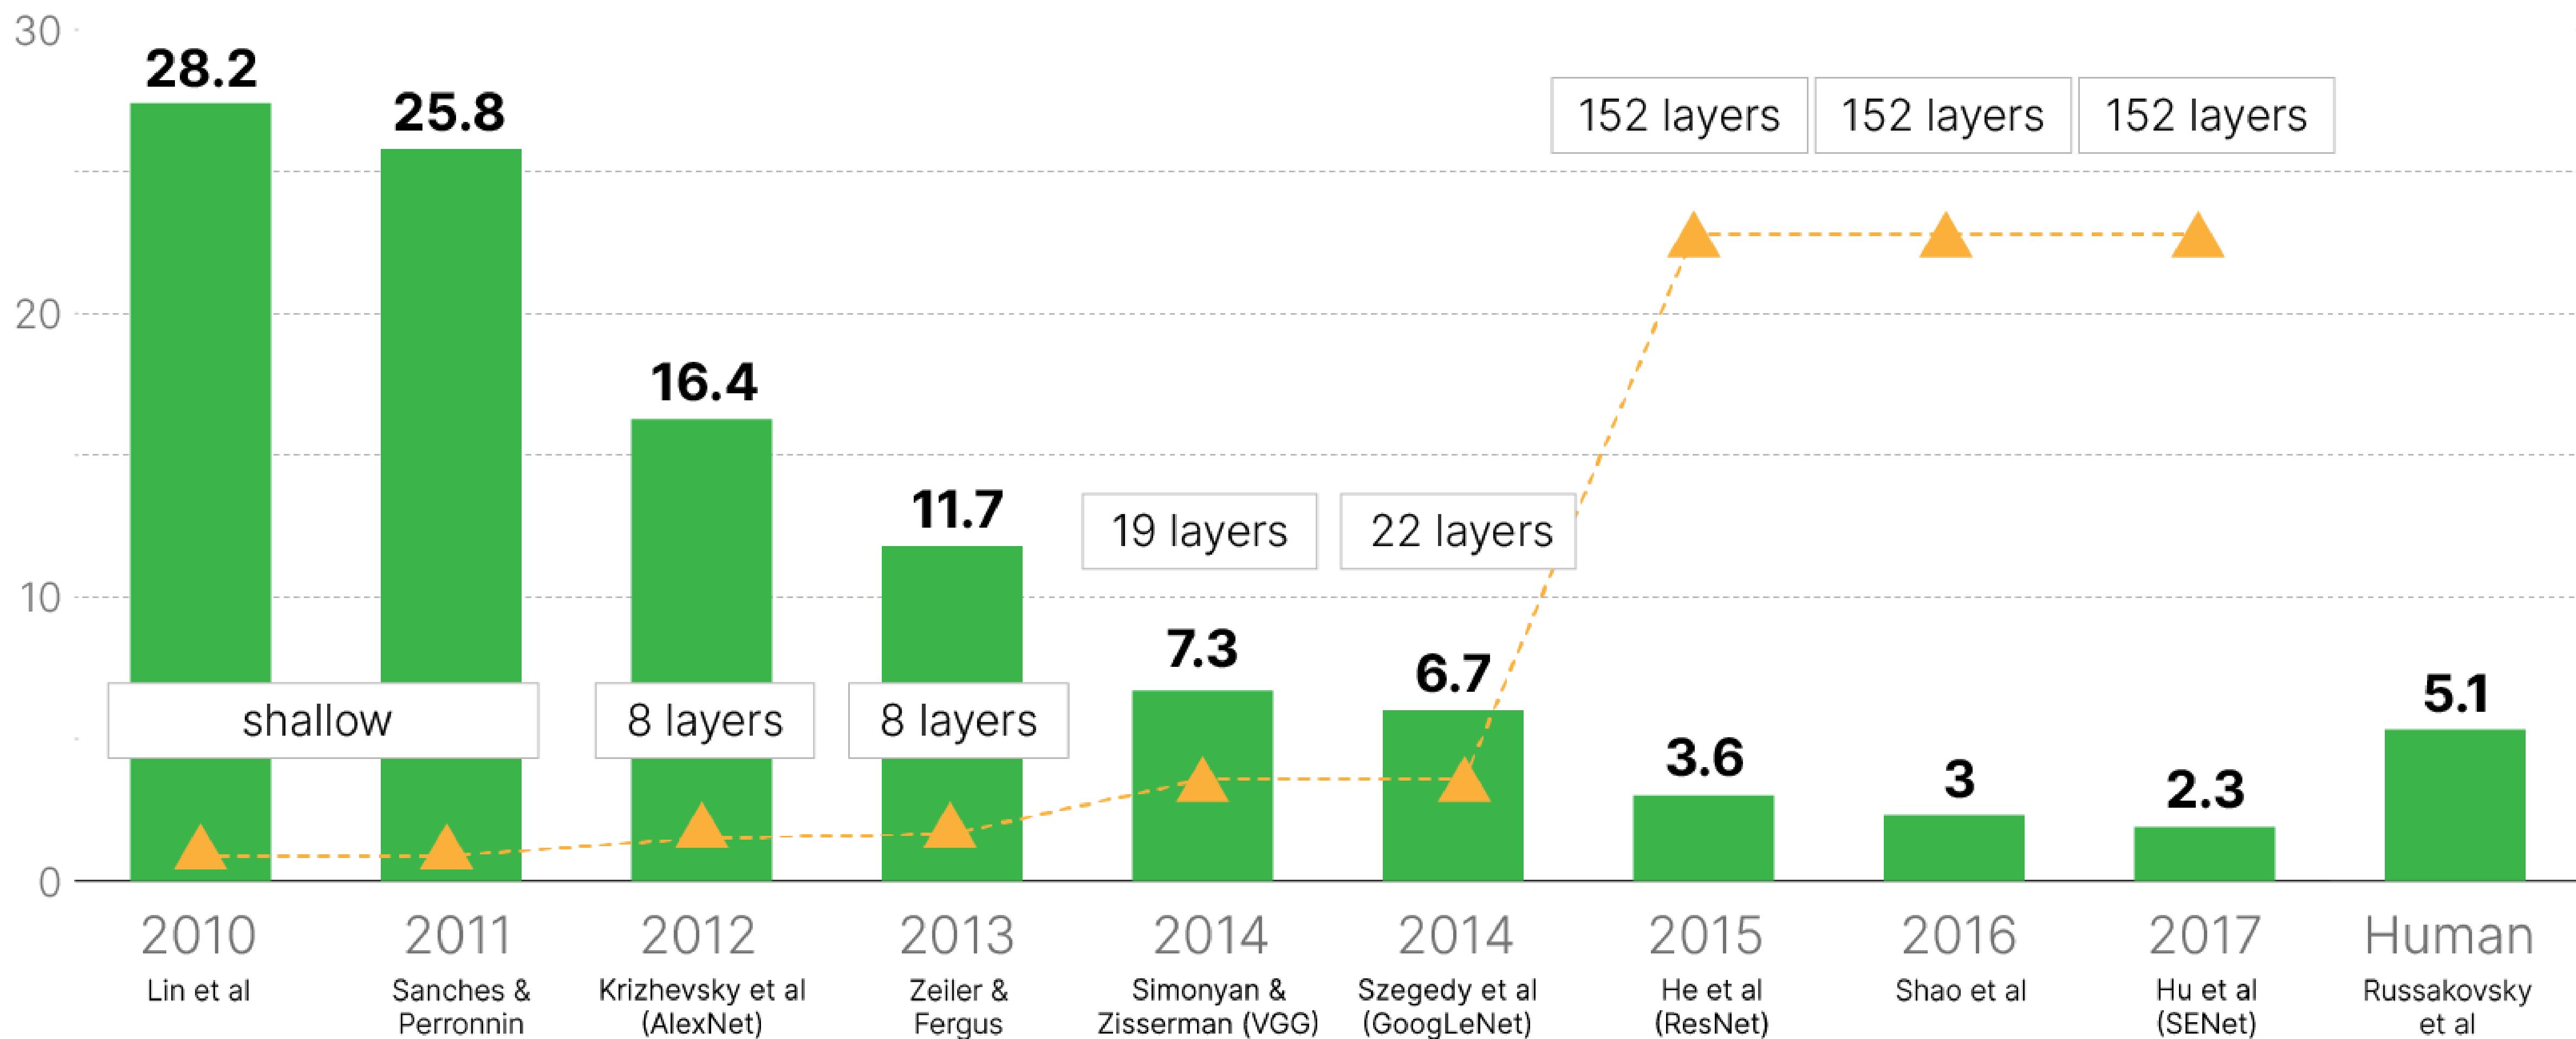
Daniil Arapov

Completely based on msu.ai lectures

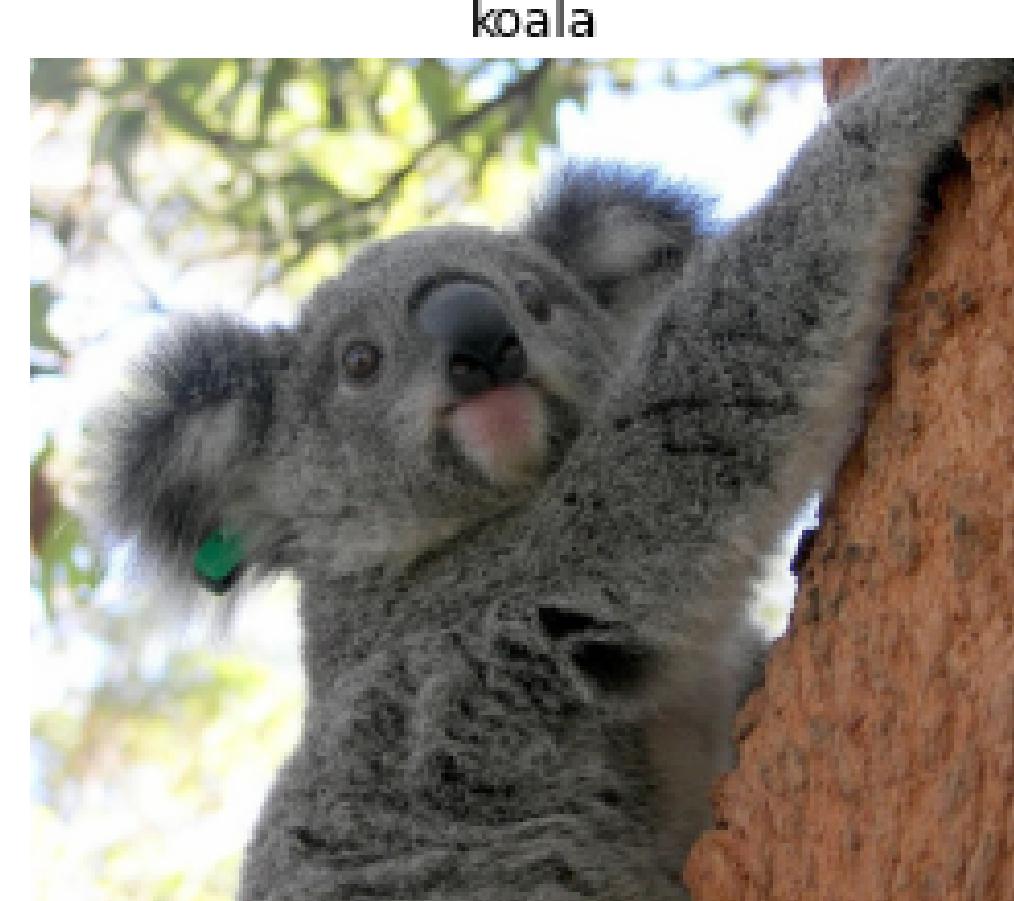
Test questions for today (up to 1 point):

1. What is your name?
2. What is TA's name?
3. How you think, what graded activity will give you the second usual test point?

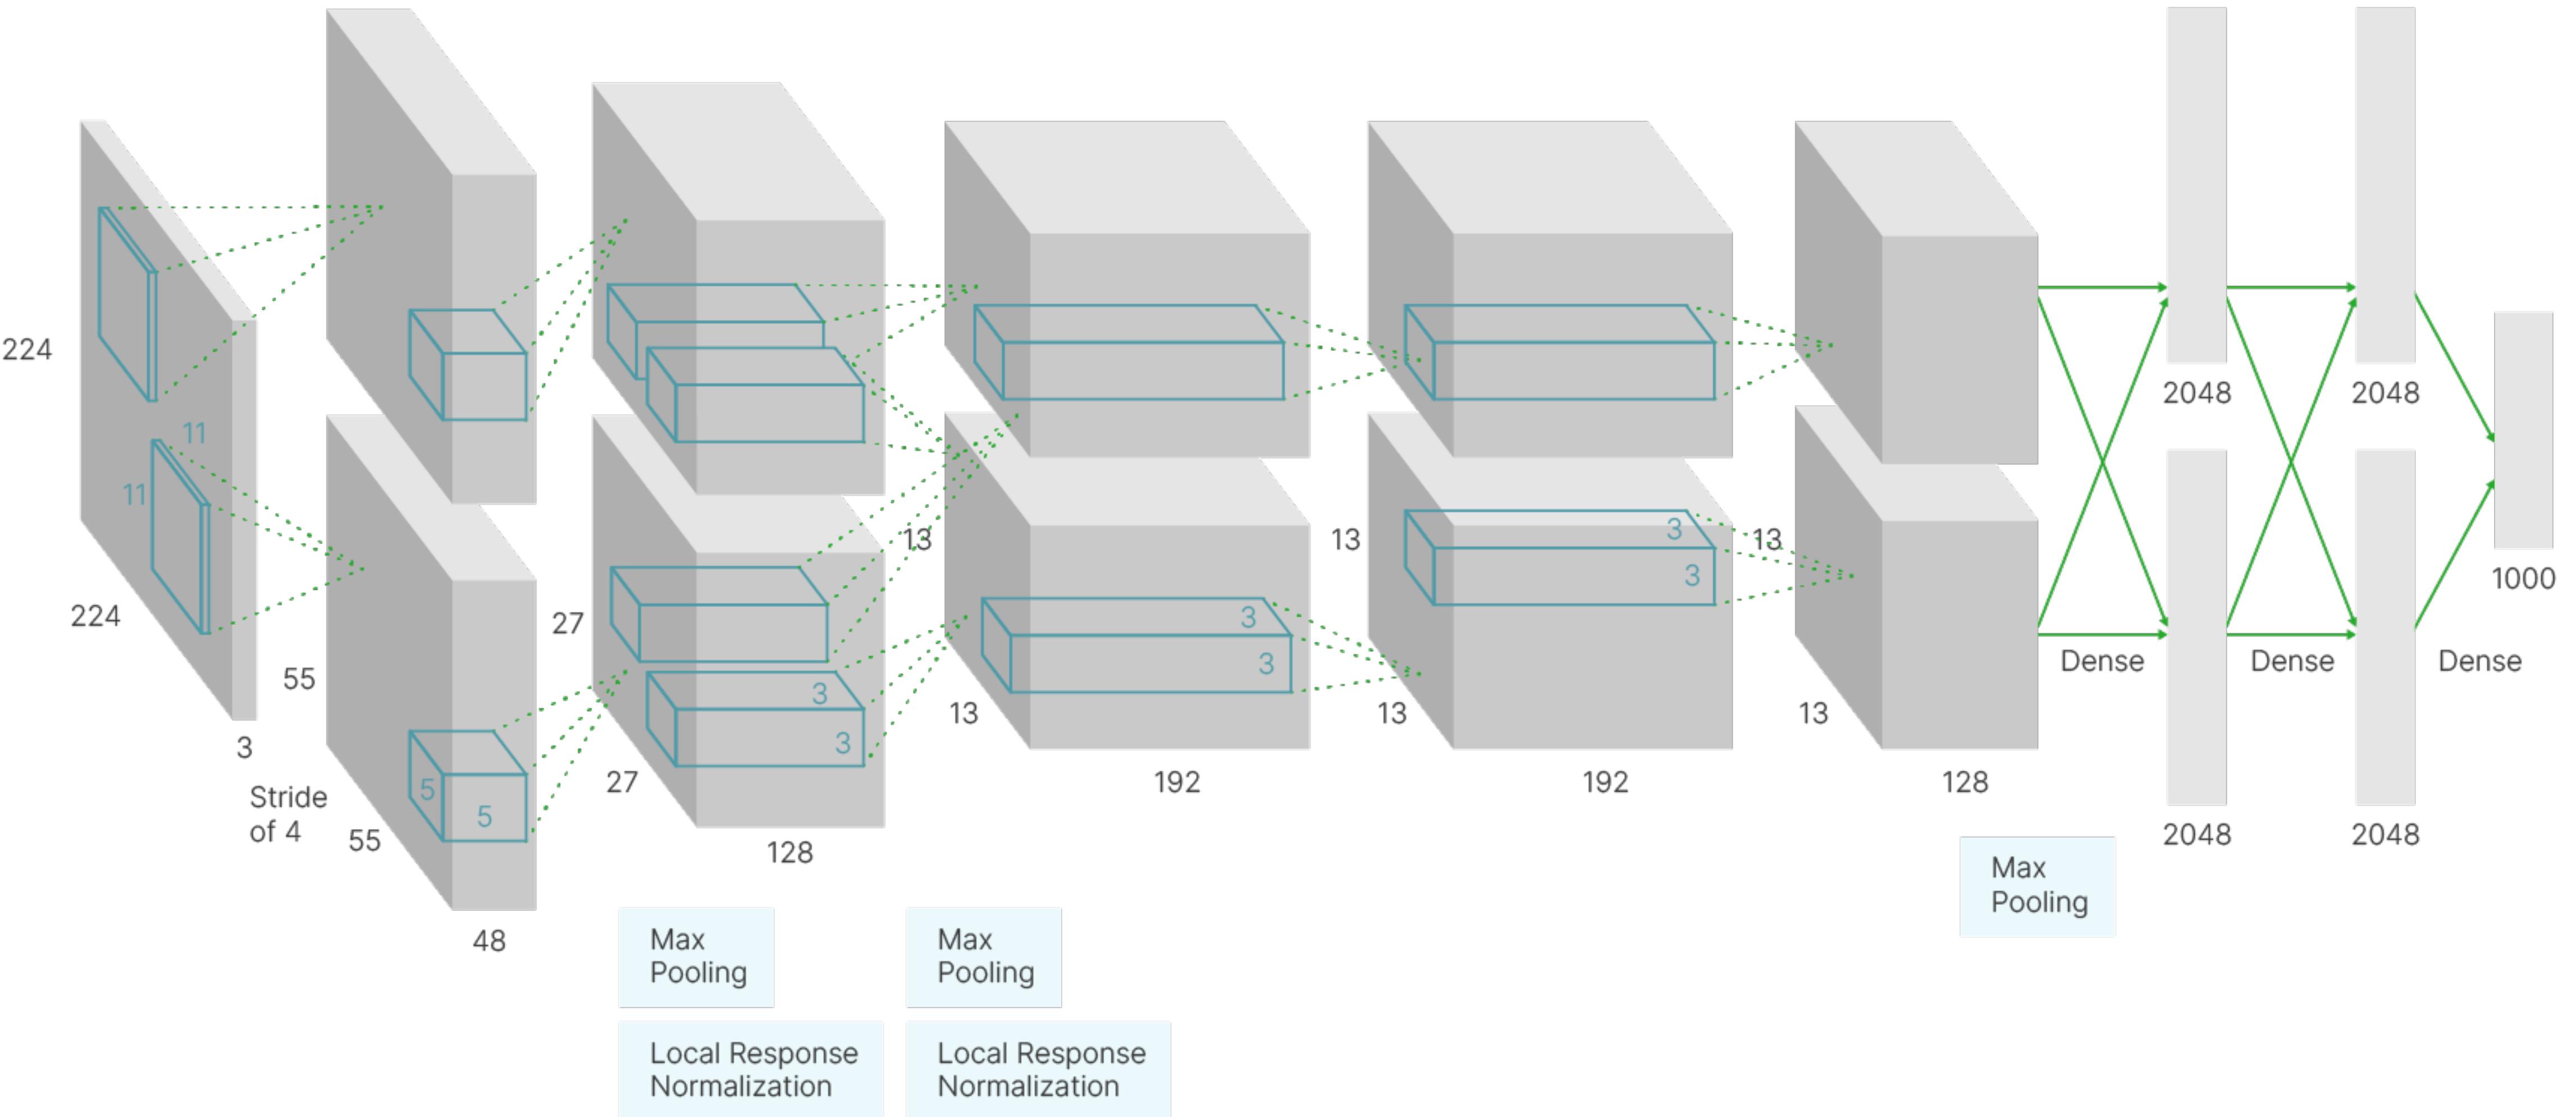
ImageNet Large Scale Visual Recognition Challenge



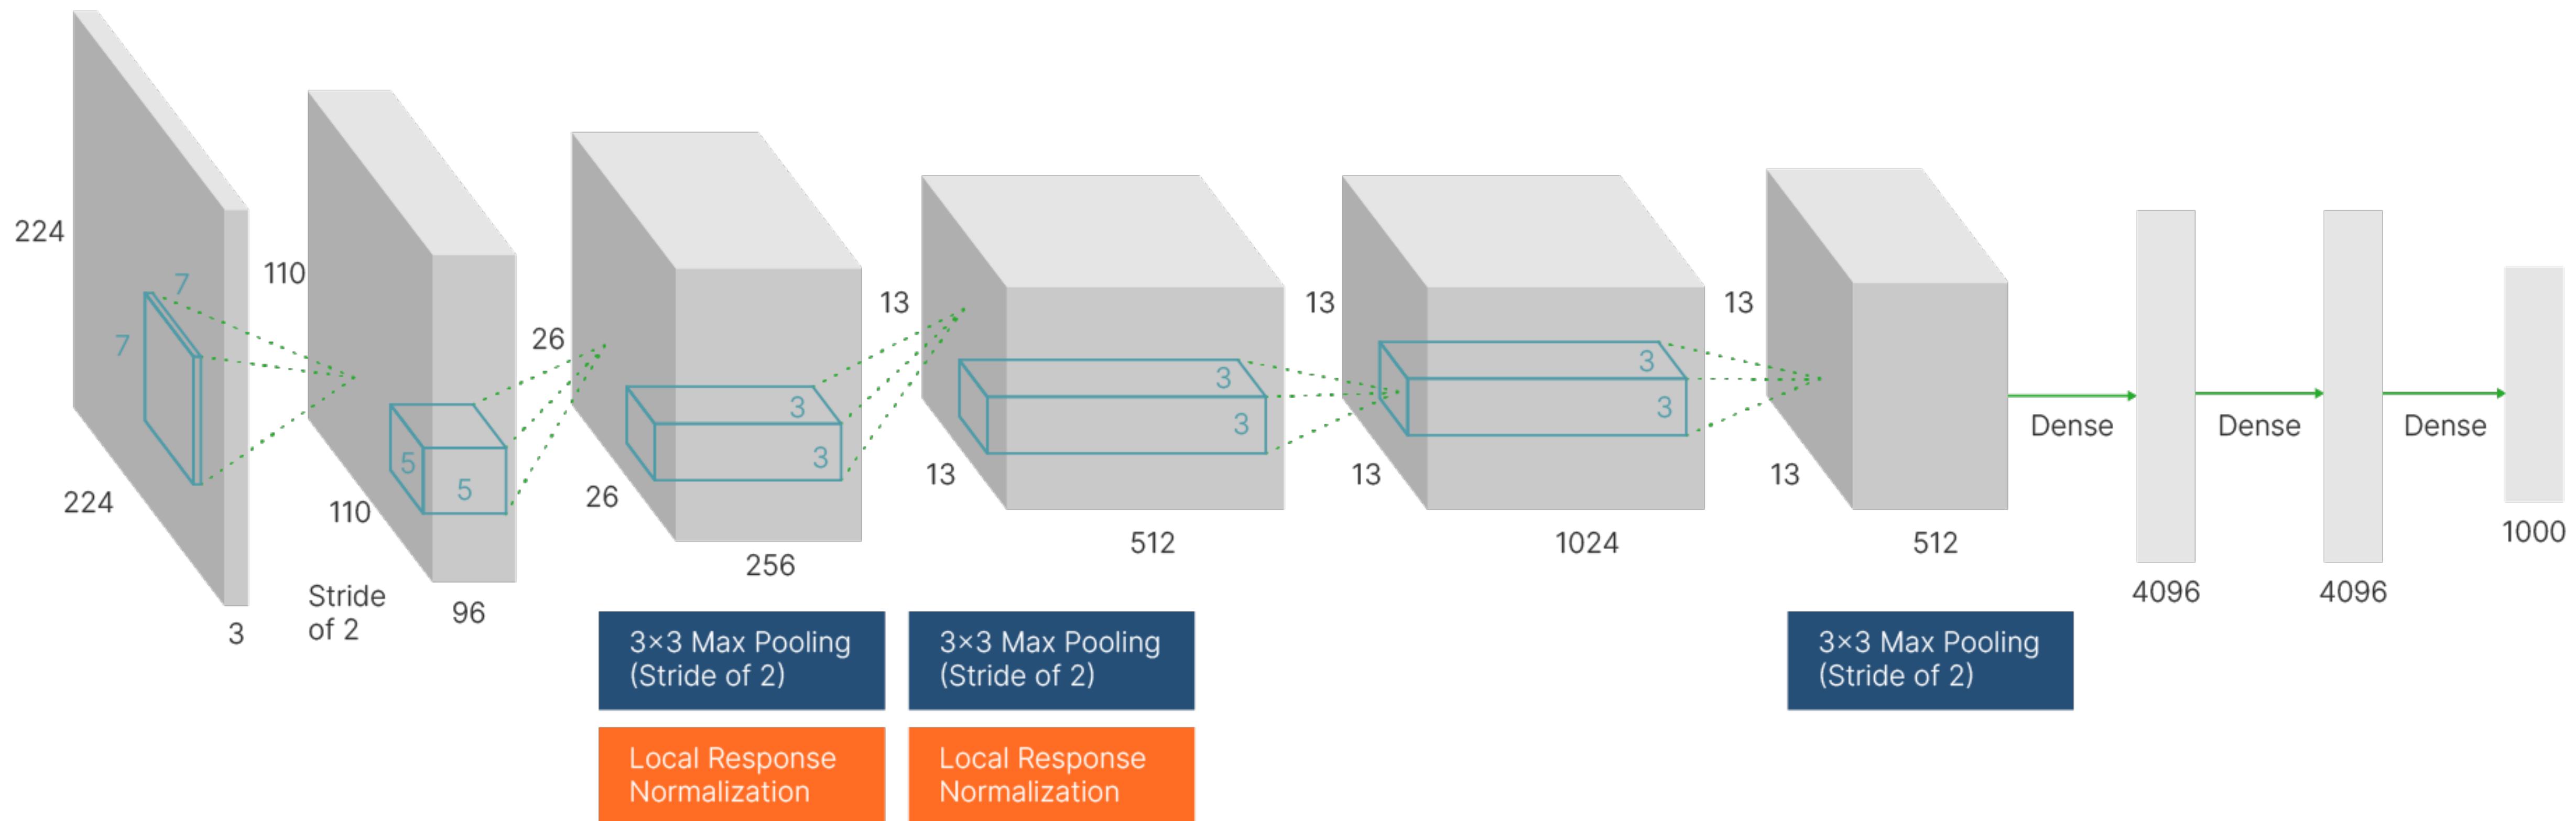
ImageNet Large Scale Visual Recognition Challenge (2)



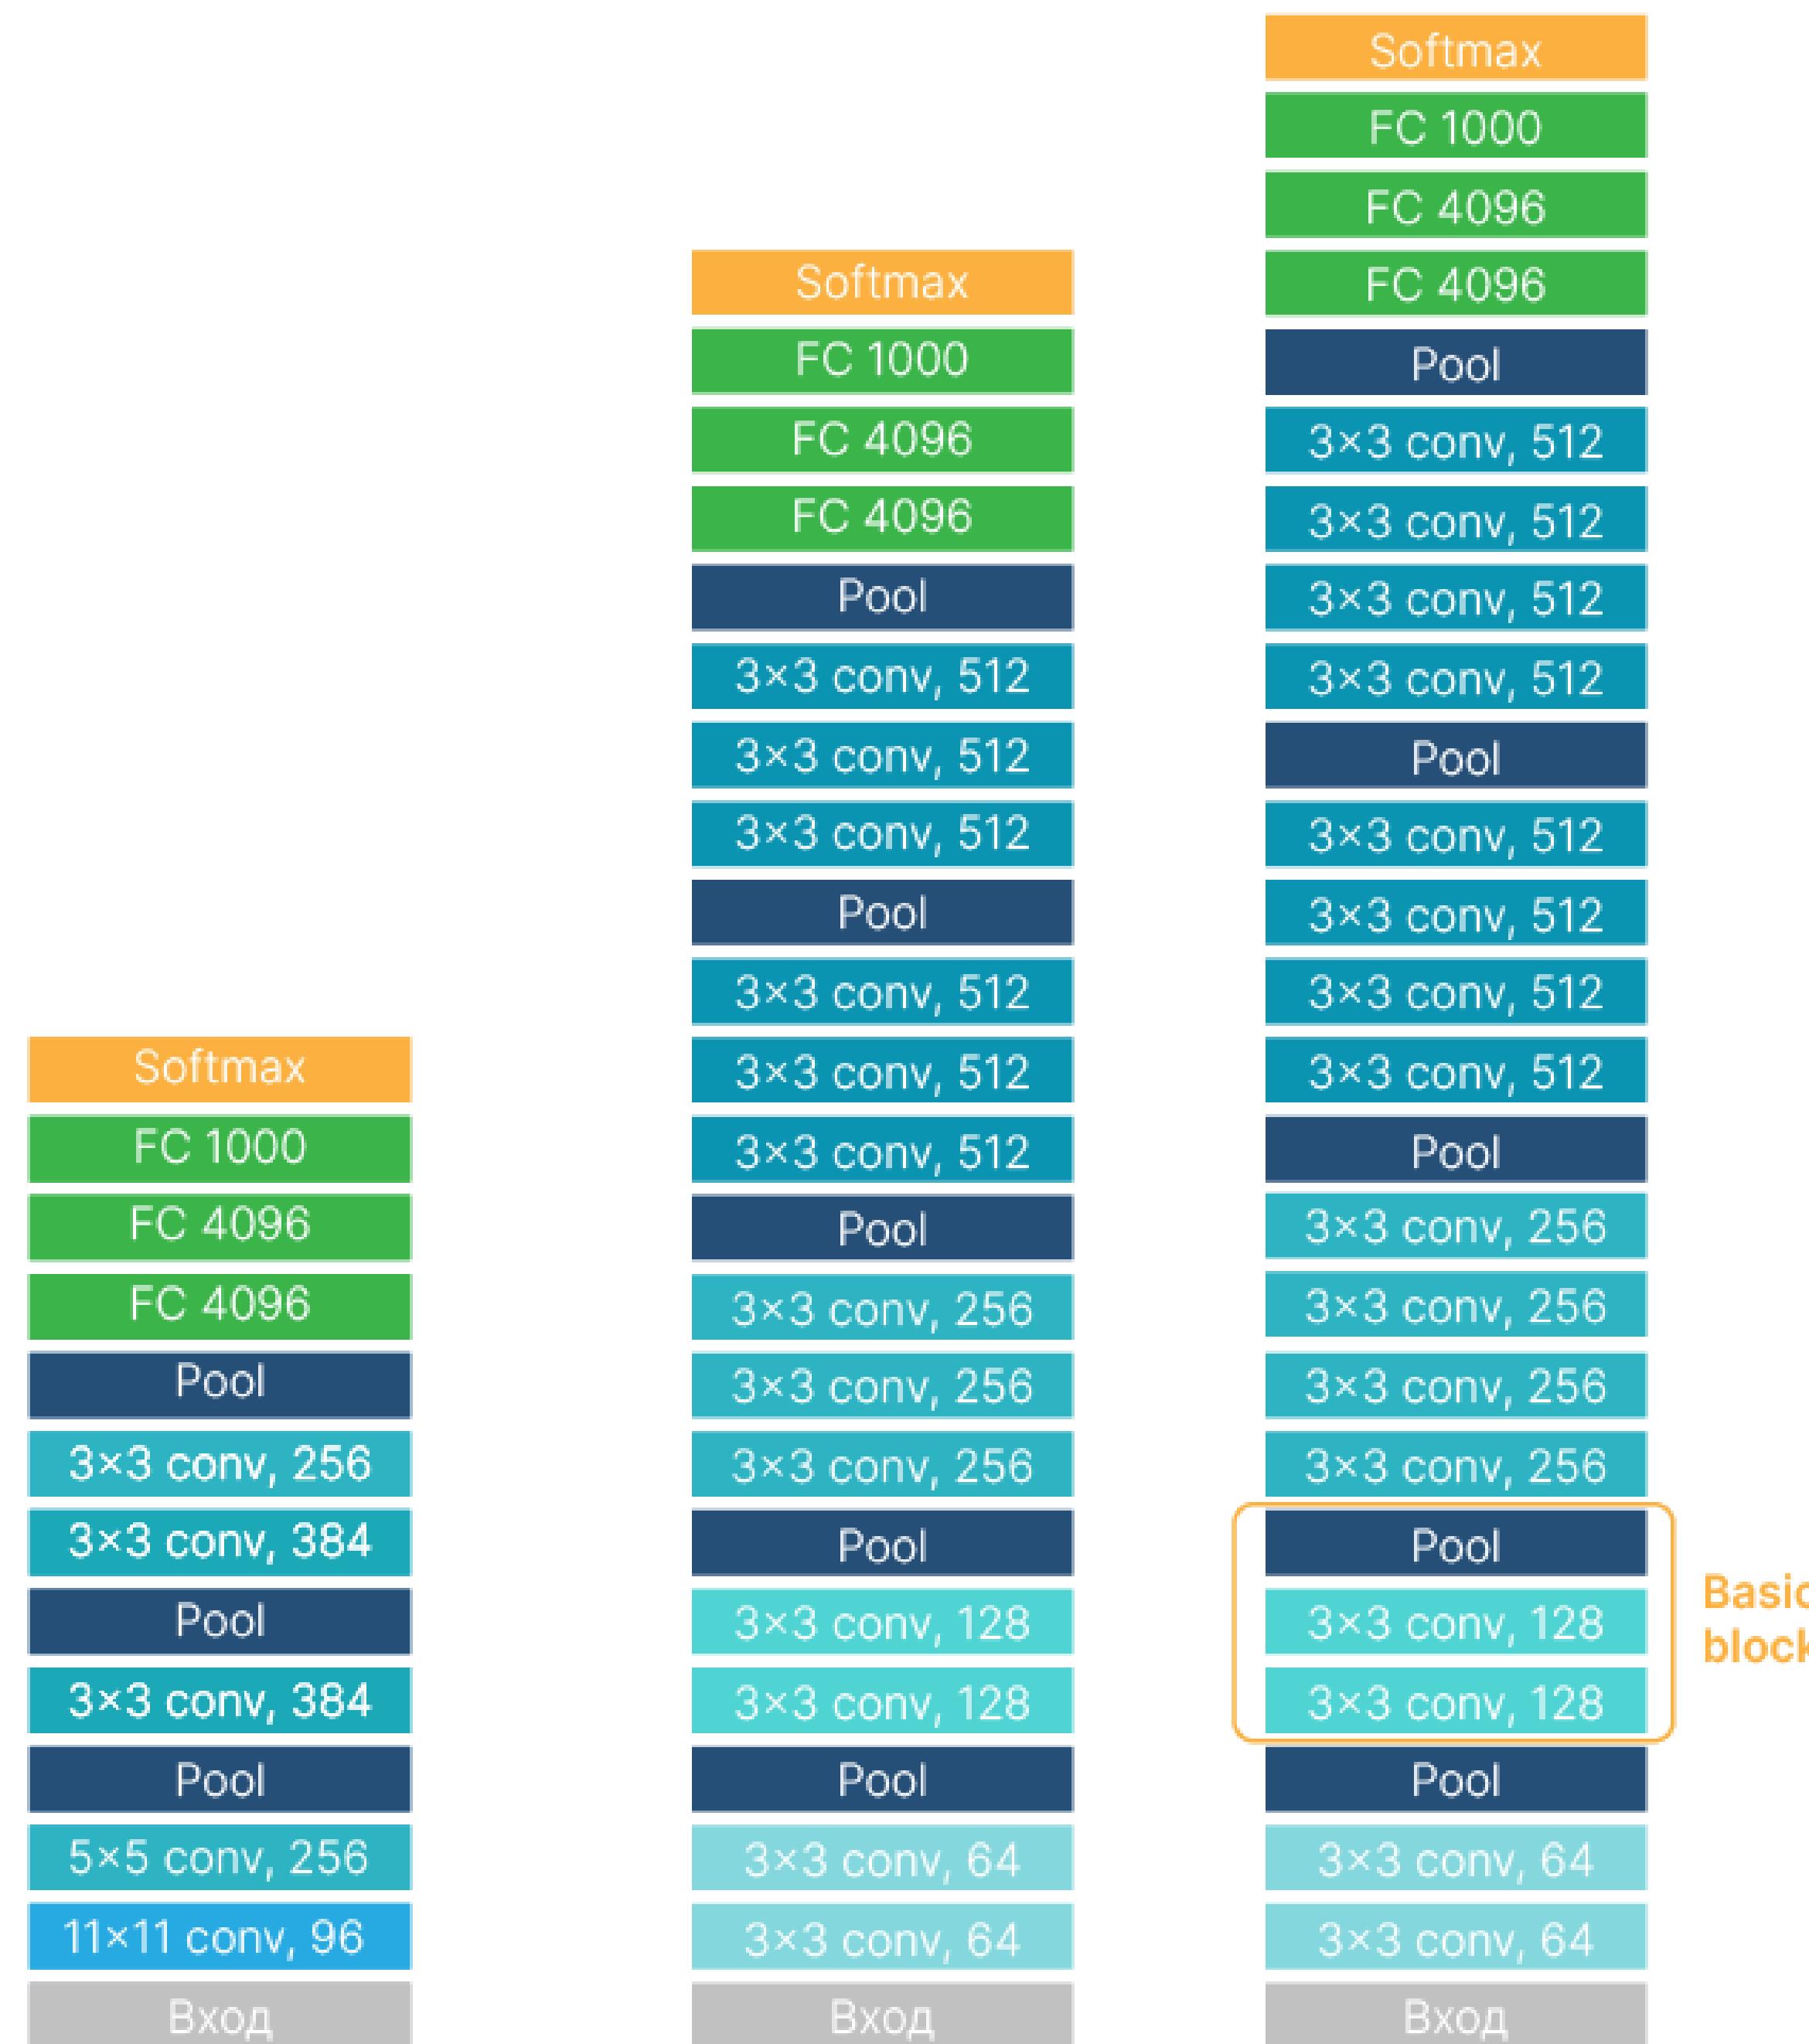
AlexNet (2012)



ZFNet (2013)



VGGNet (2014)

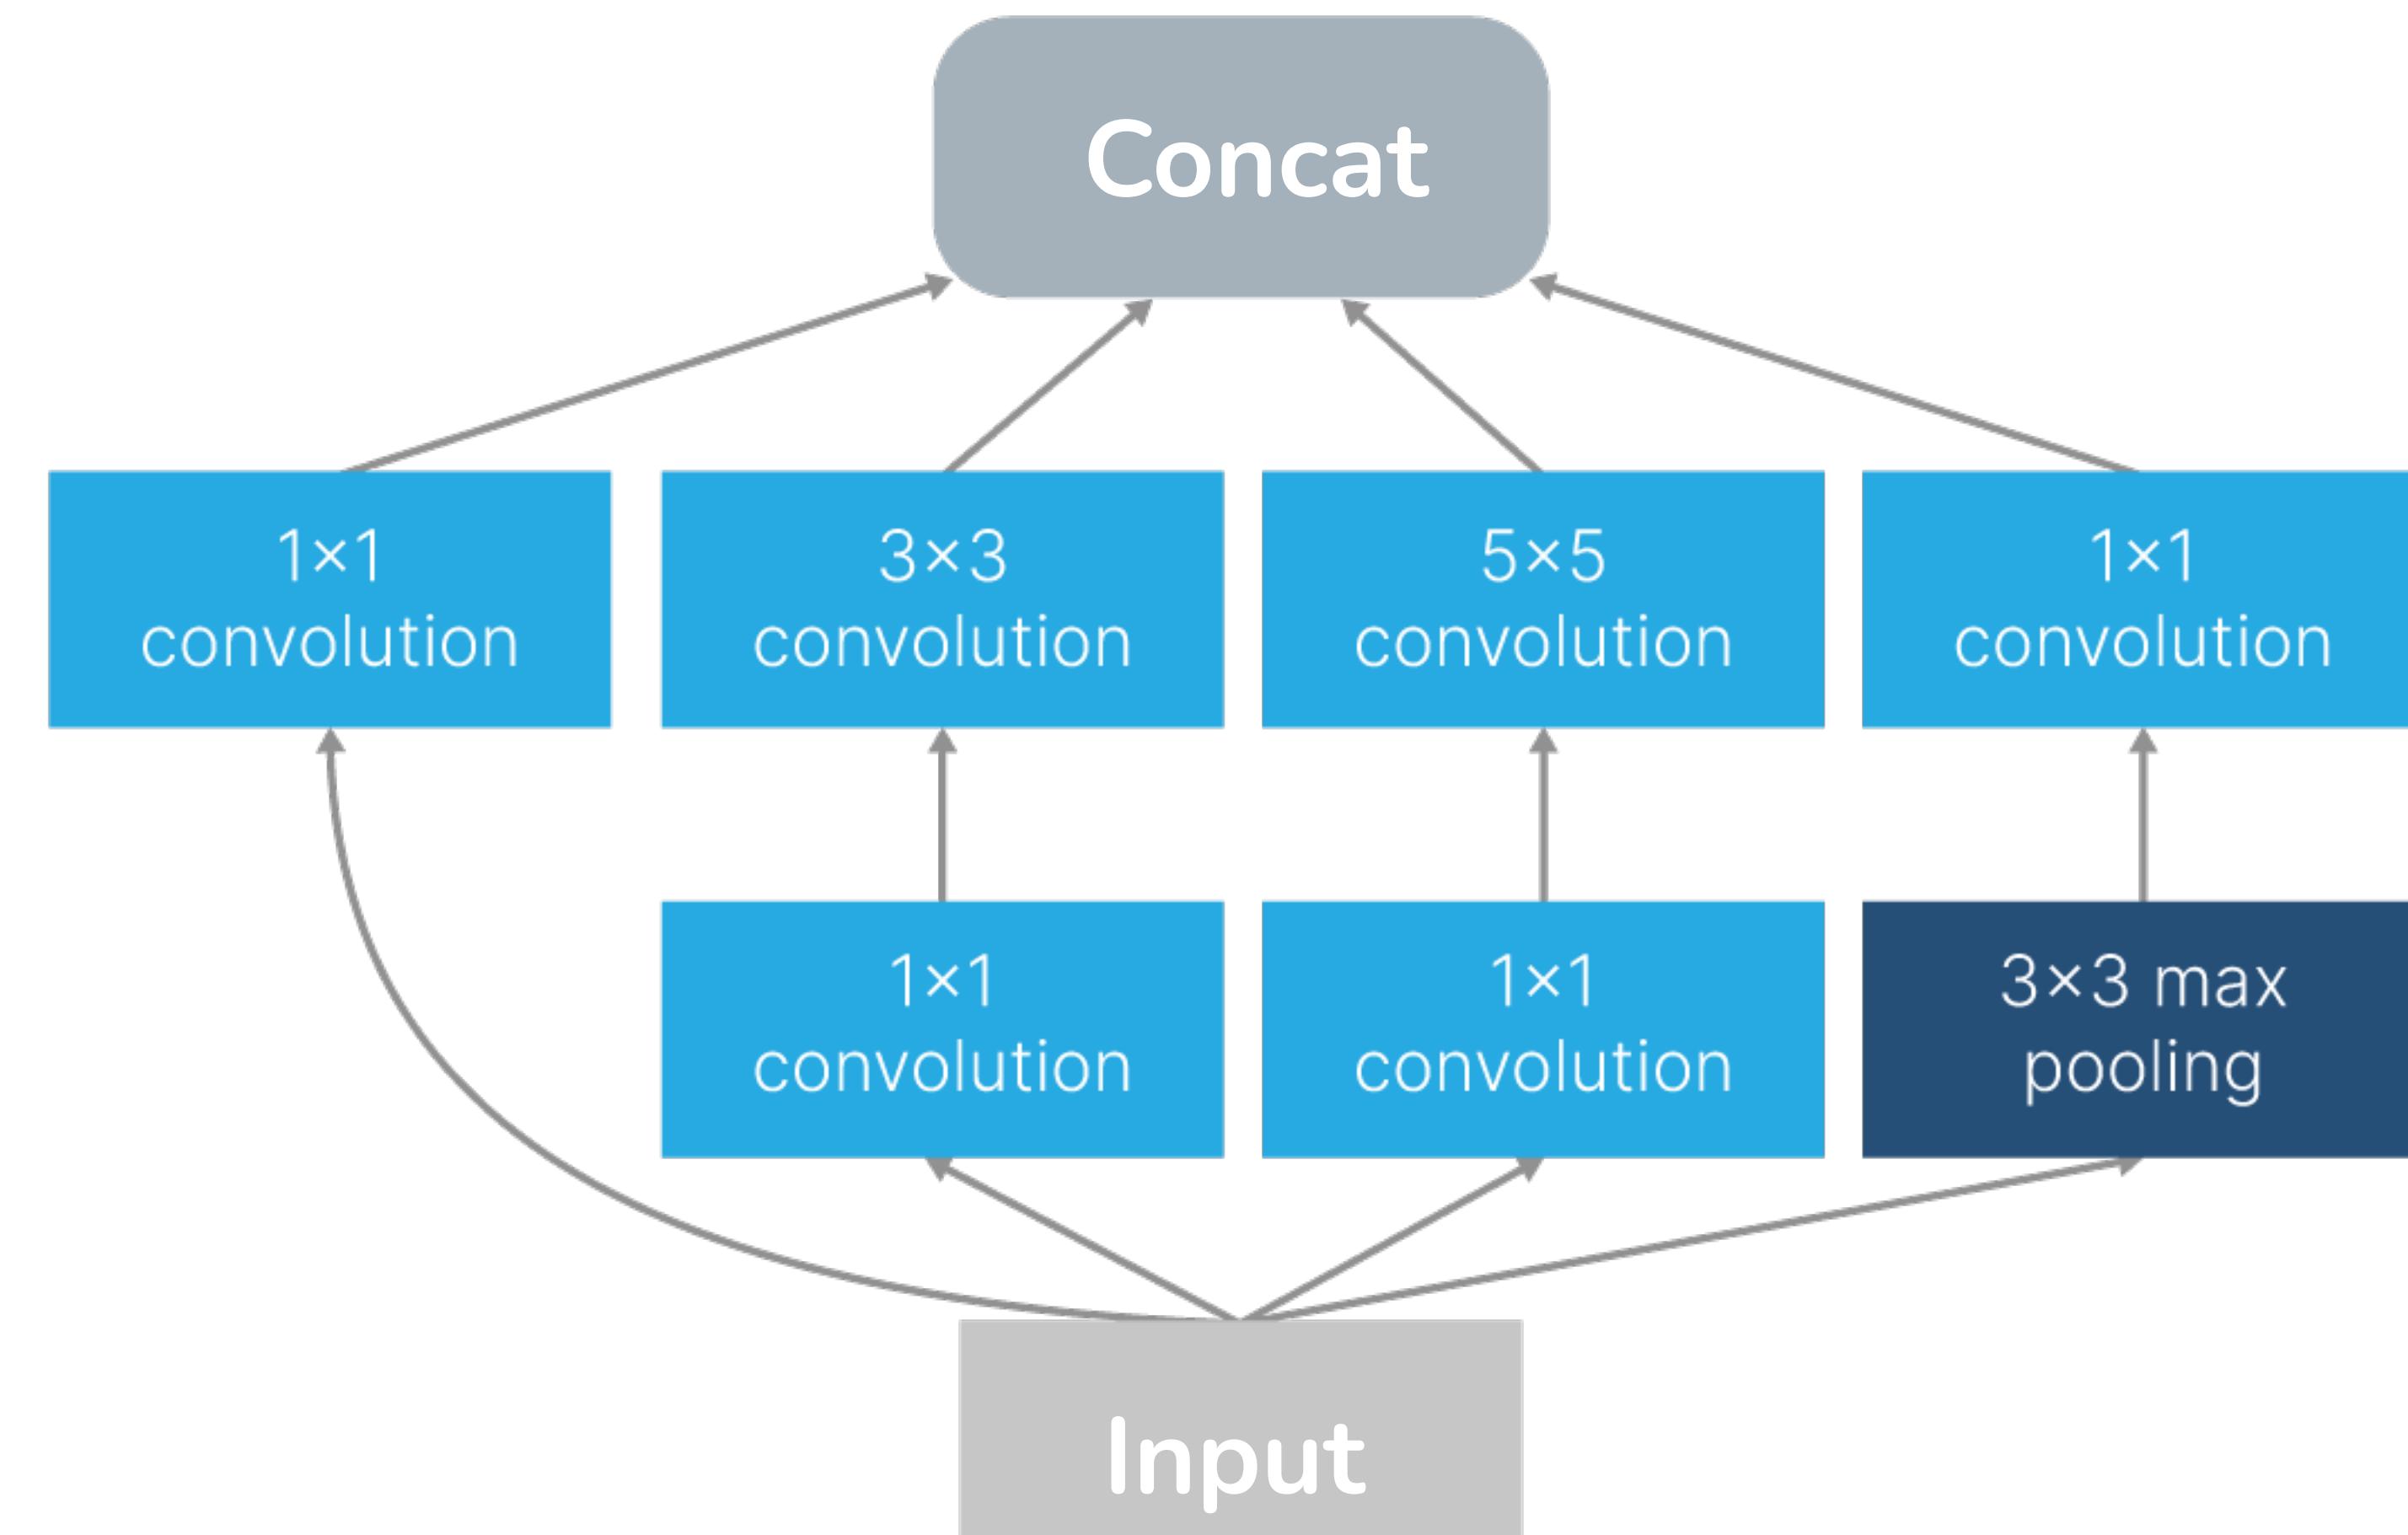


AlexNet

VGG16

VGG19

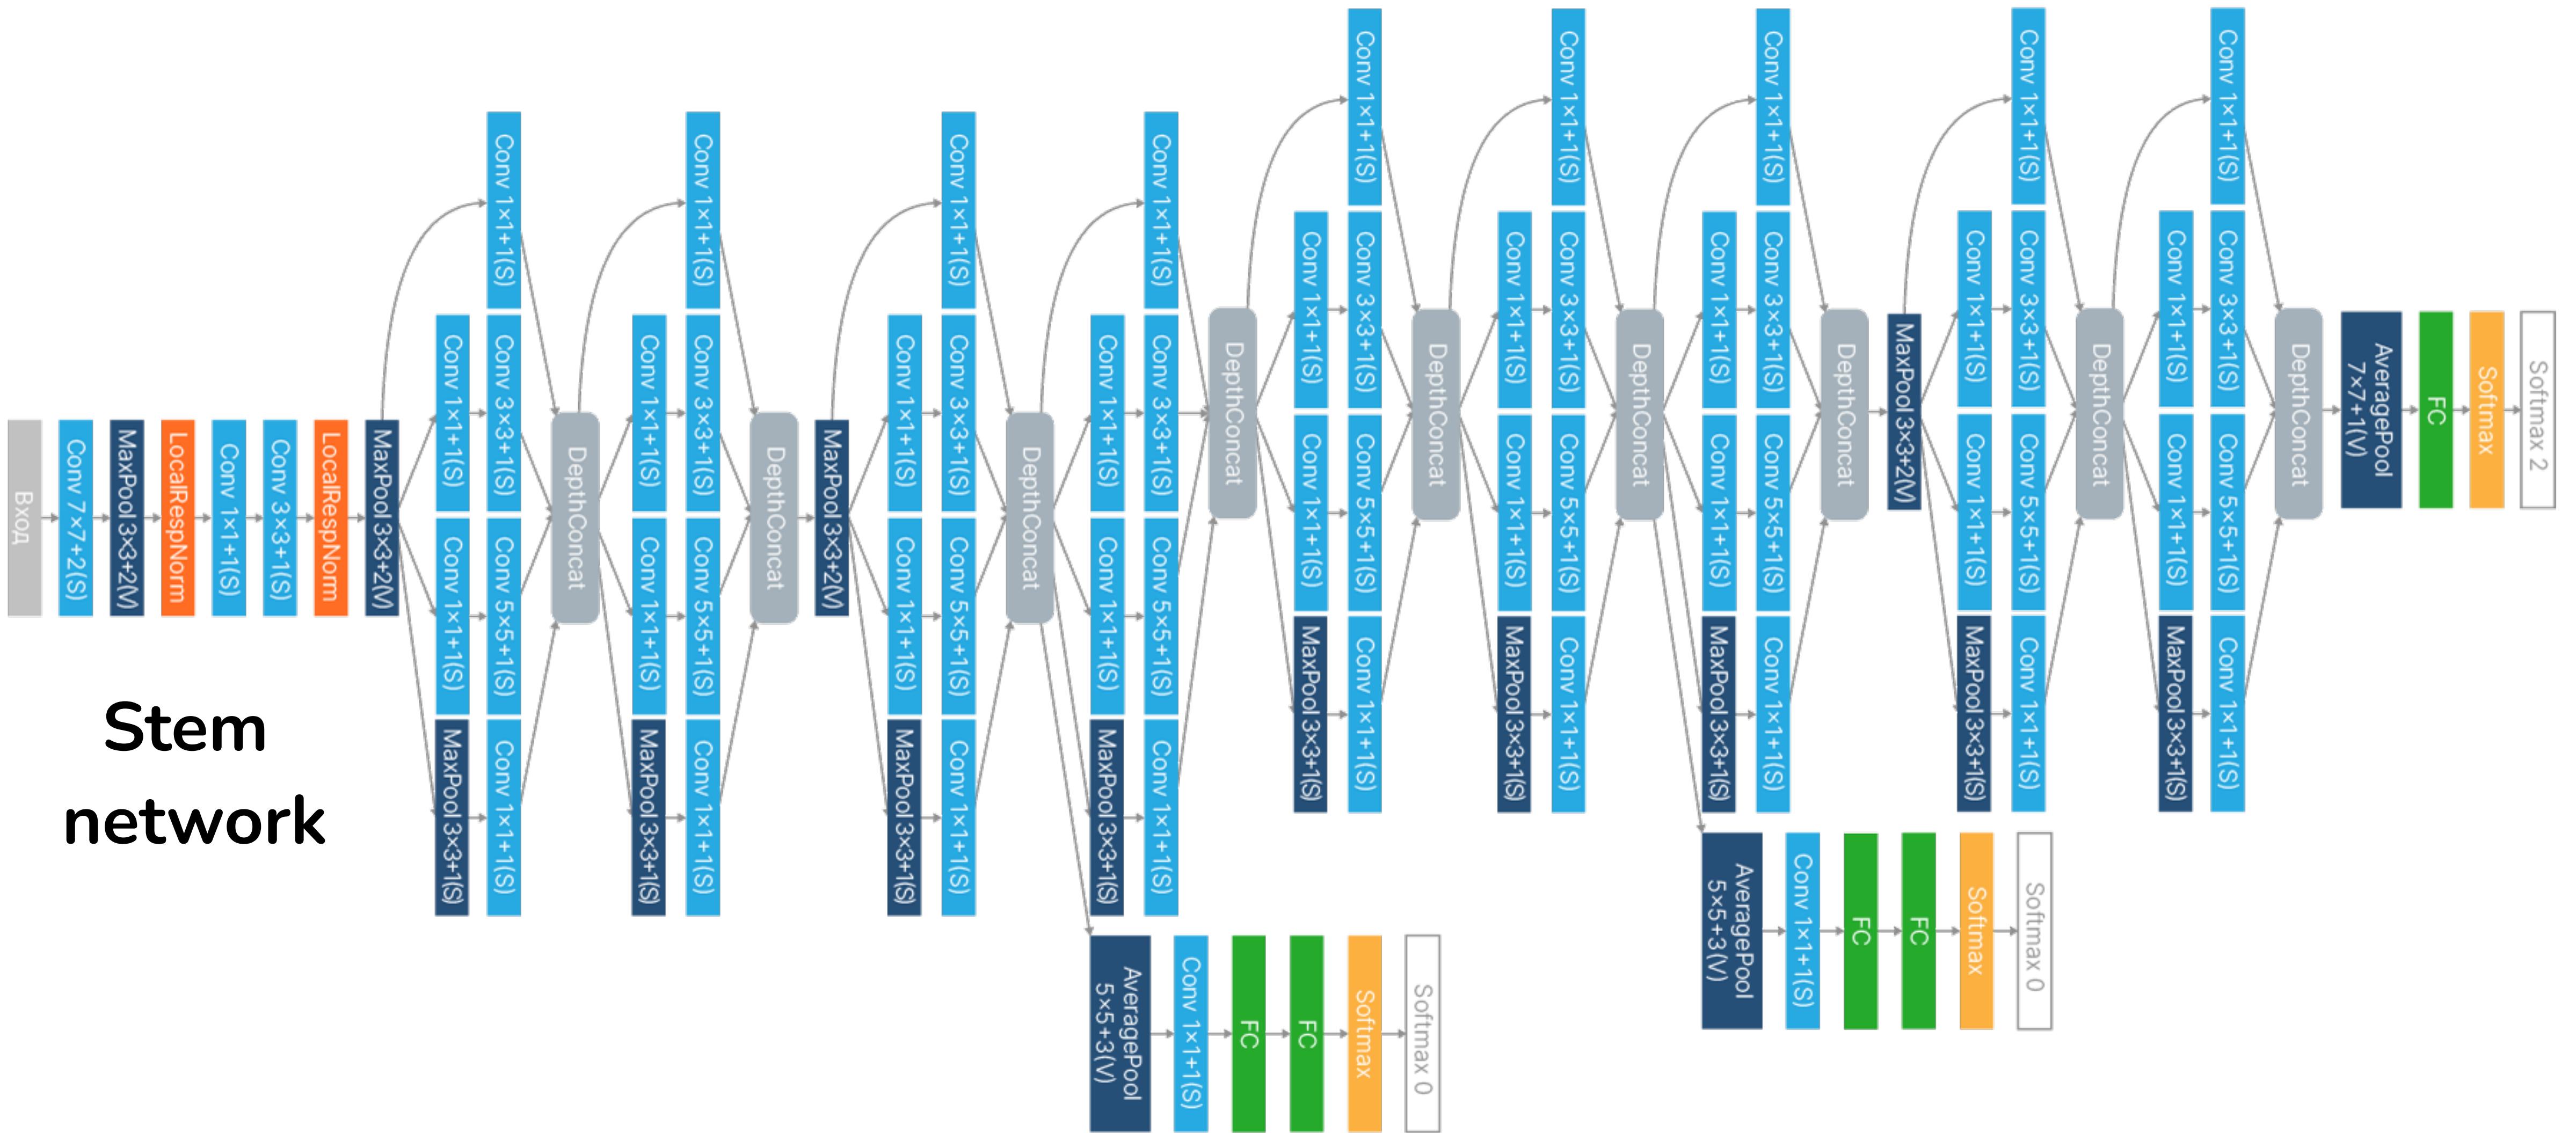
GoogLeNet (2014) | Inception module



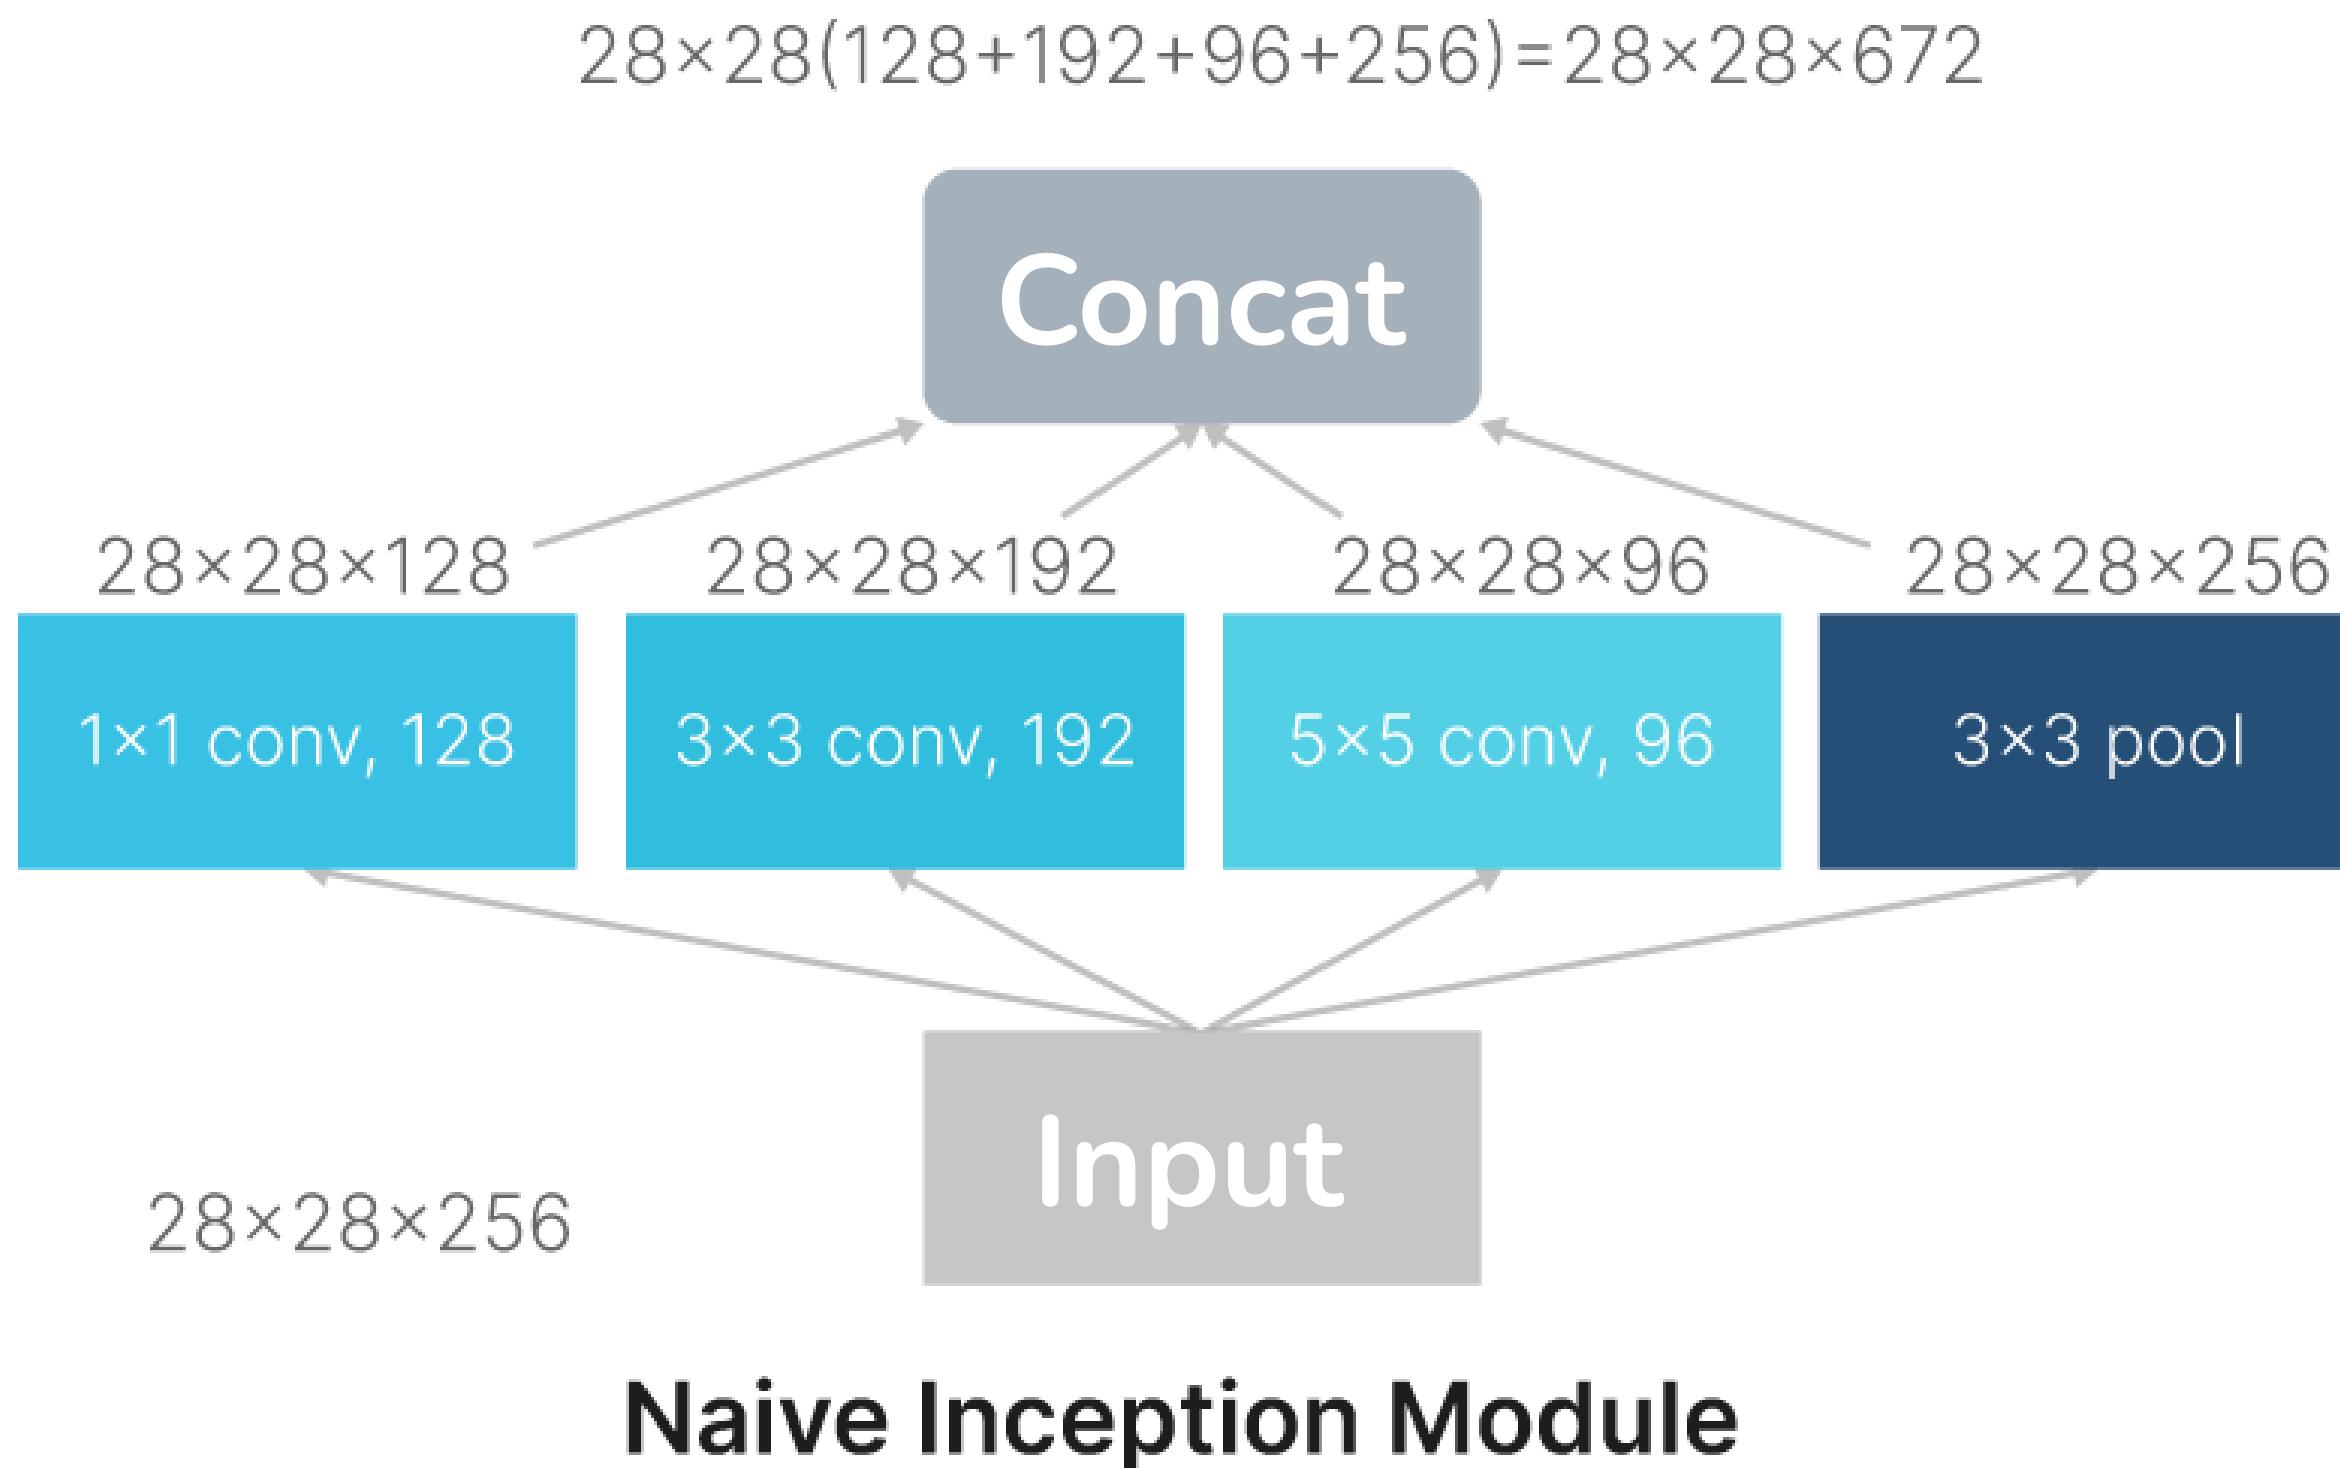
Inception Module

GoogLeNet (2014) | Architecture

Stem network



GoogLeNet (2014) | Naïve Inception Module



Conv Ops:

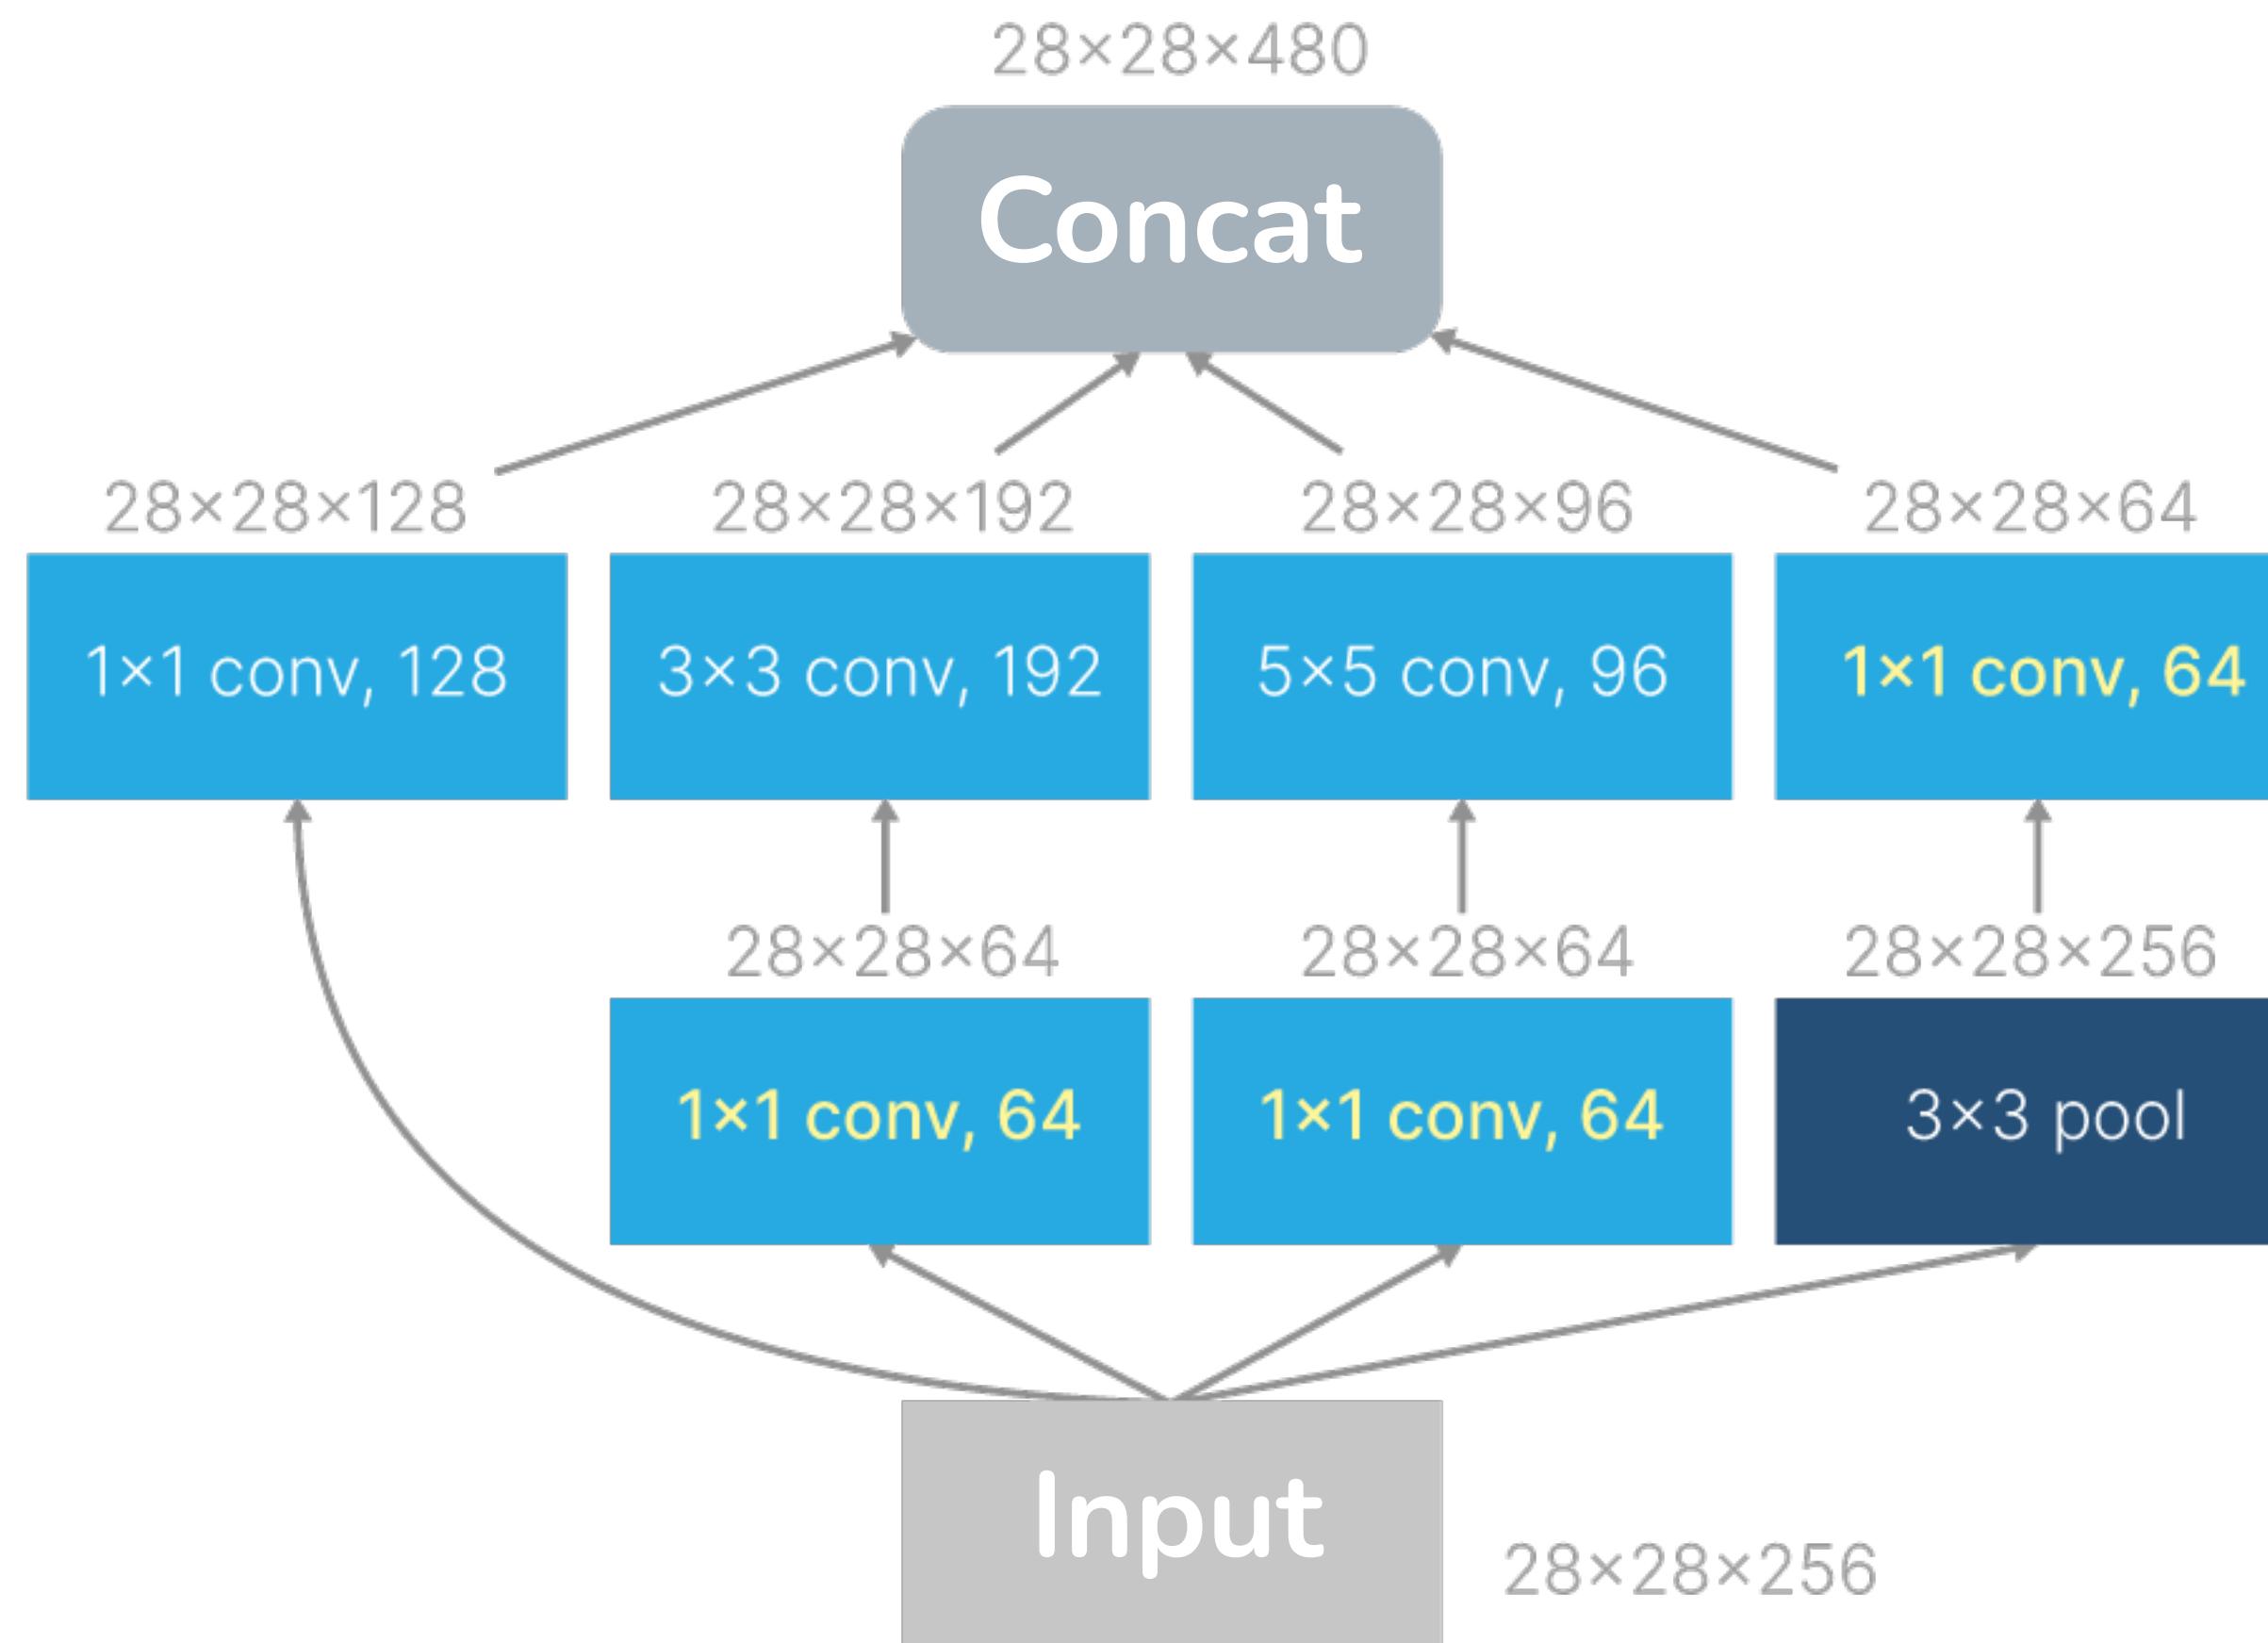
[1x1 conv, 128] $28 \times 28 \times 128 \times 1 \times 1 \times 256$

[3x3 conv, 192] $28 \times 28 \times 192 \times 3 \times 3 \times 256$

[5x5 conv, 96] $28 \times 28 \times 96 \times 5 \times 5 \times 256$

Total: 854M ops

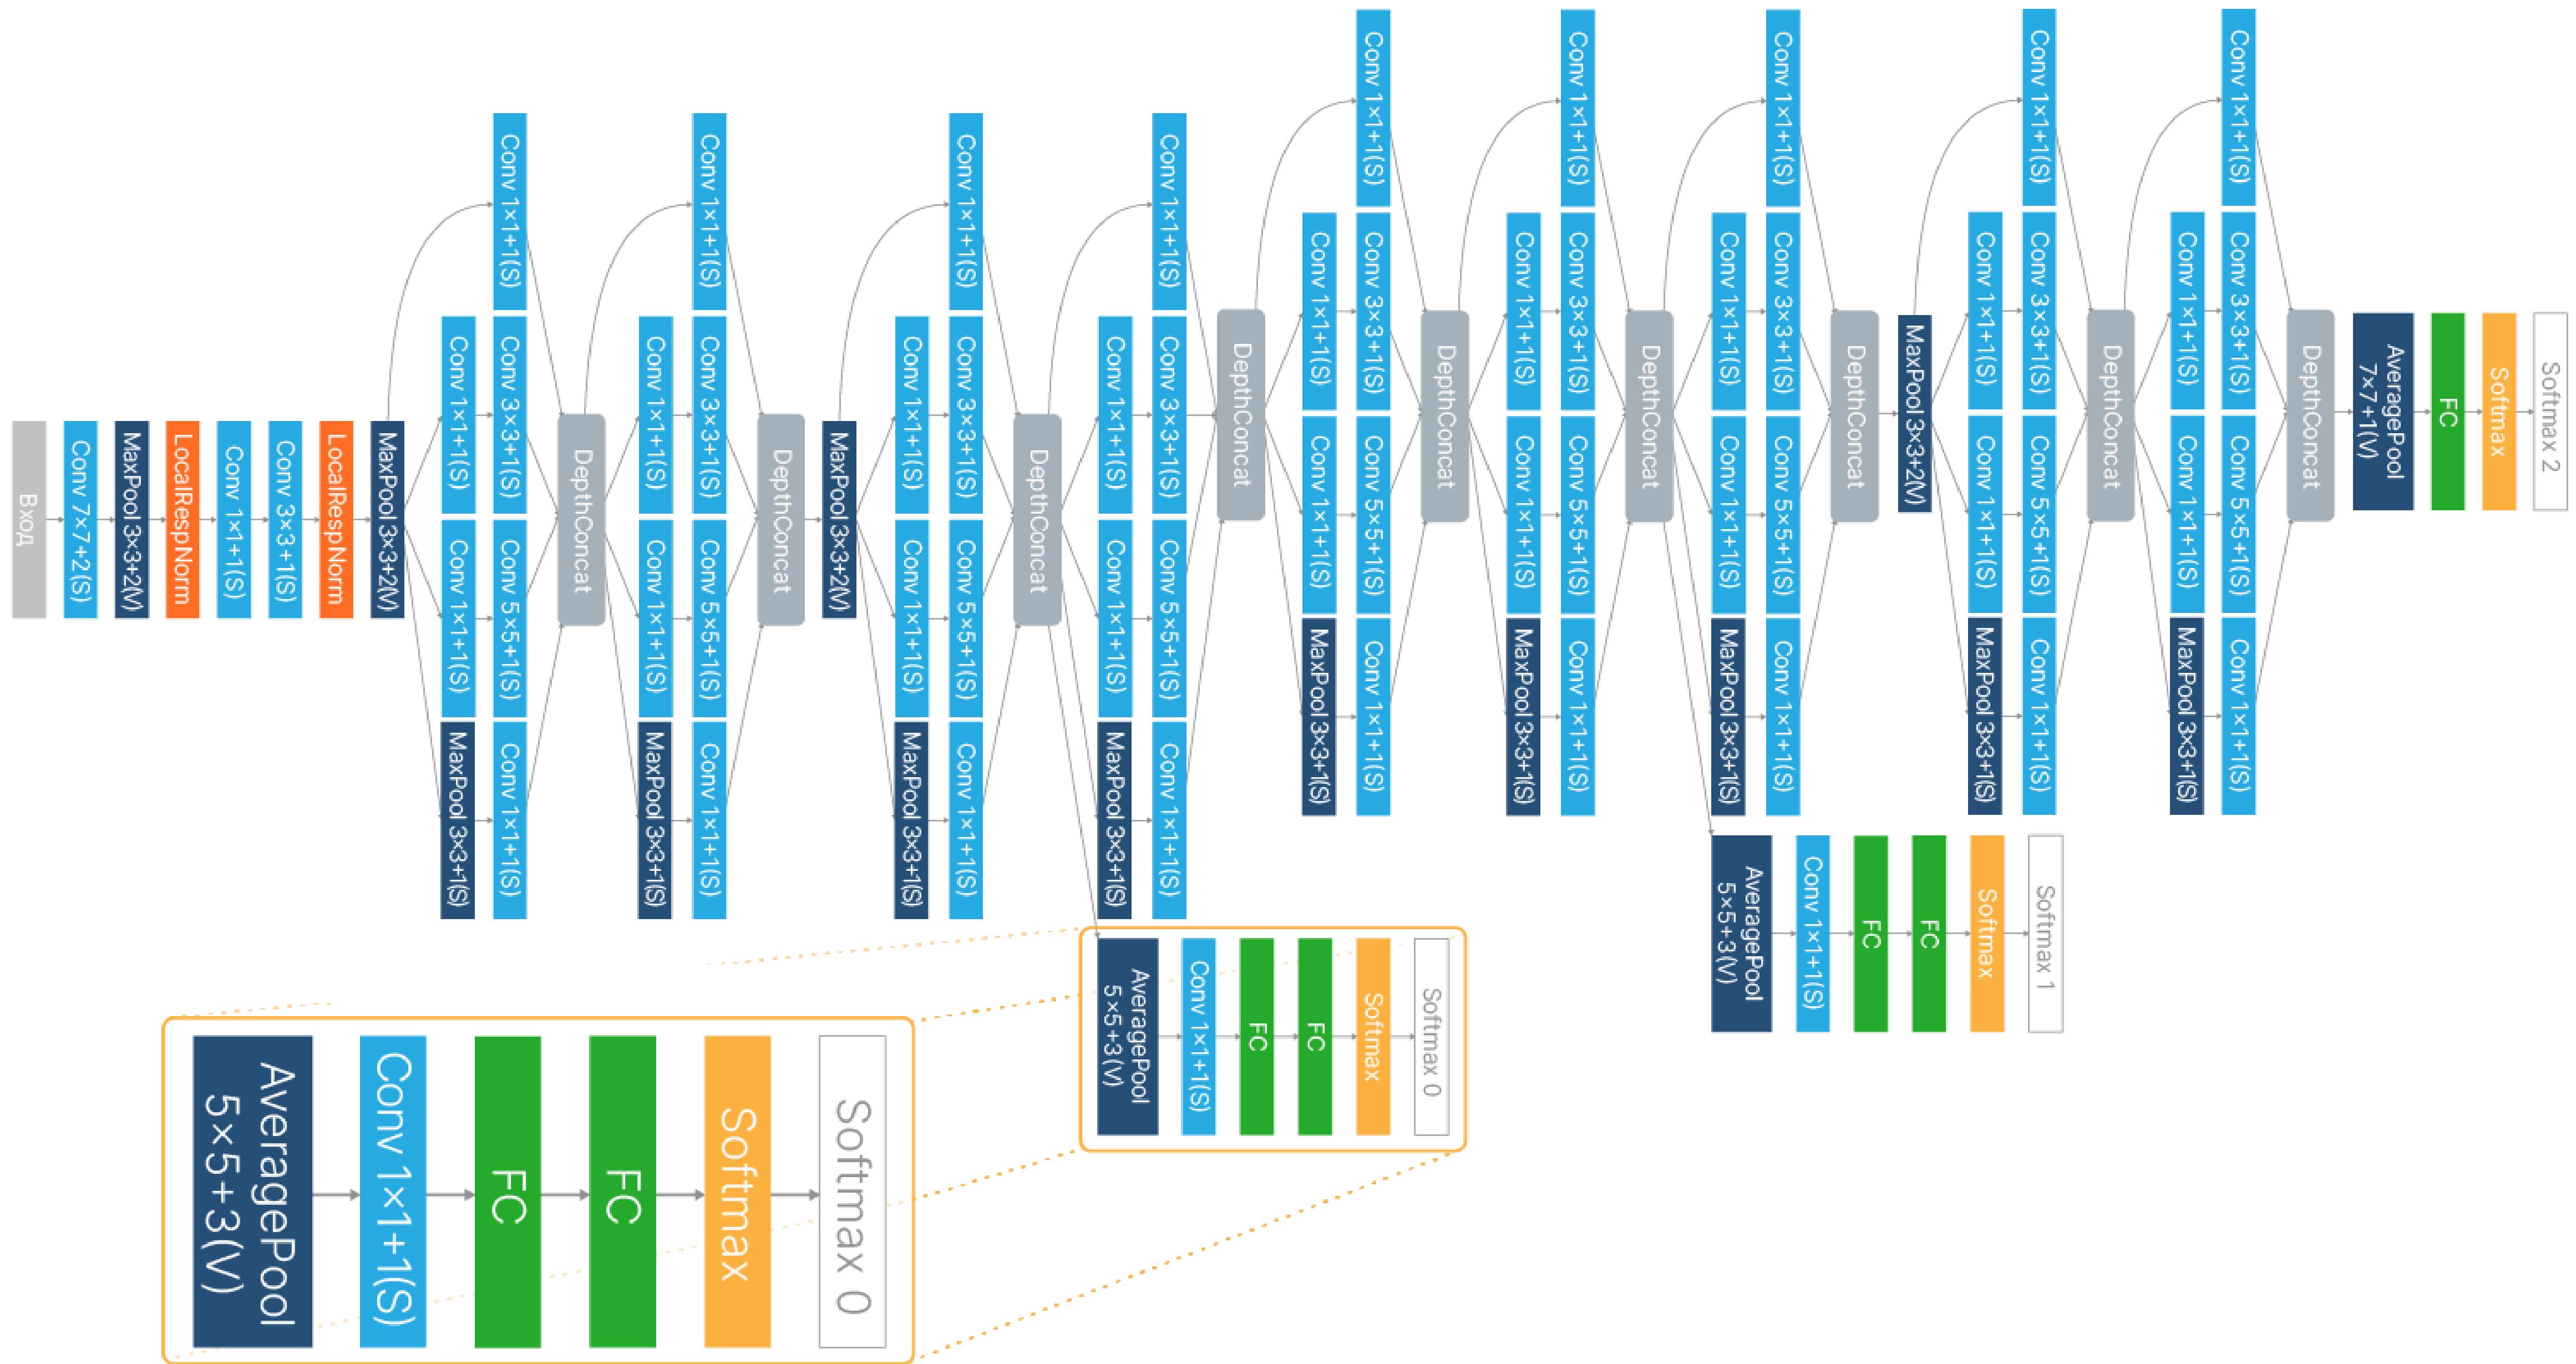
GoogLeNet (2014) | Inception Module



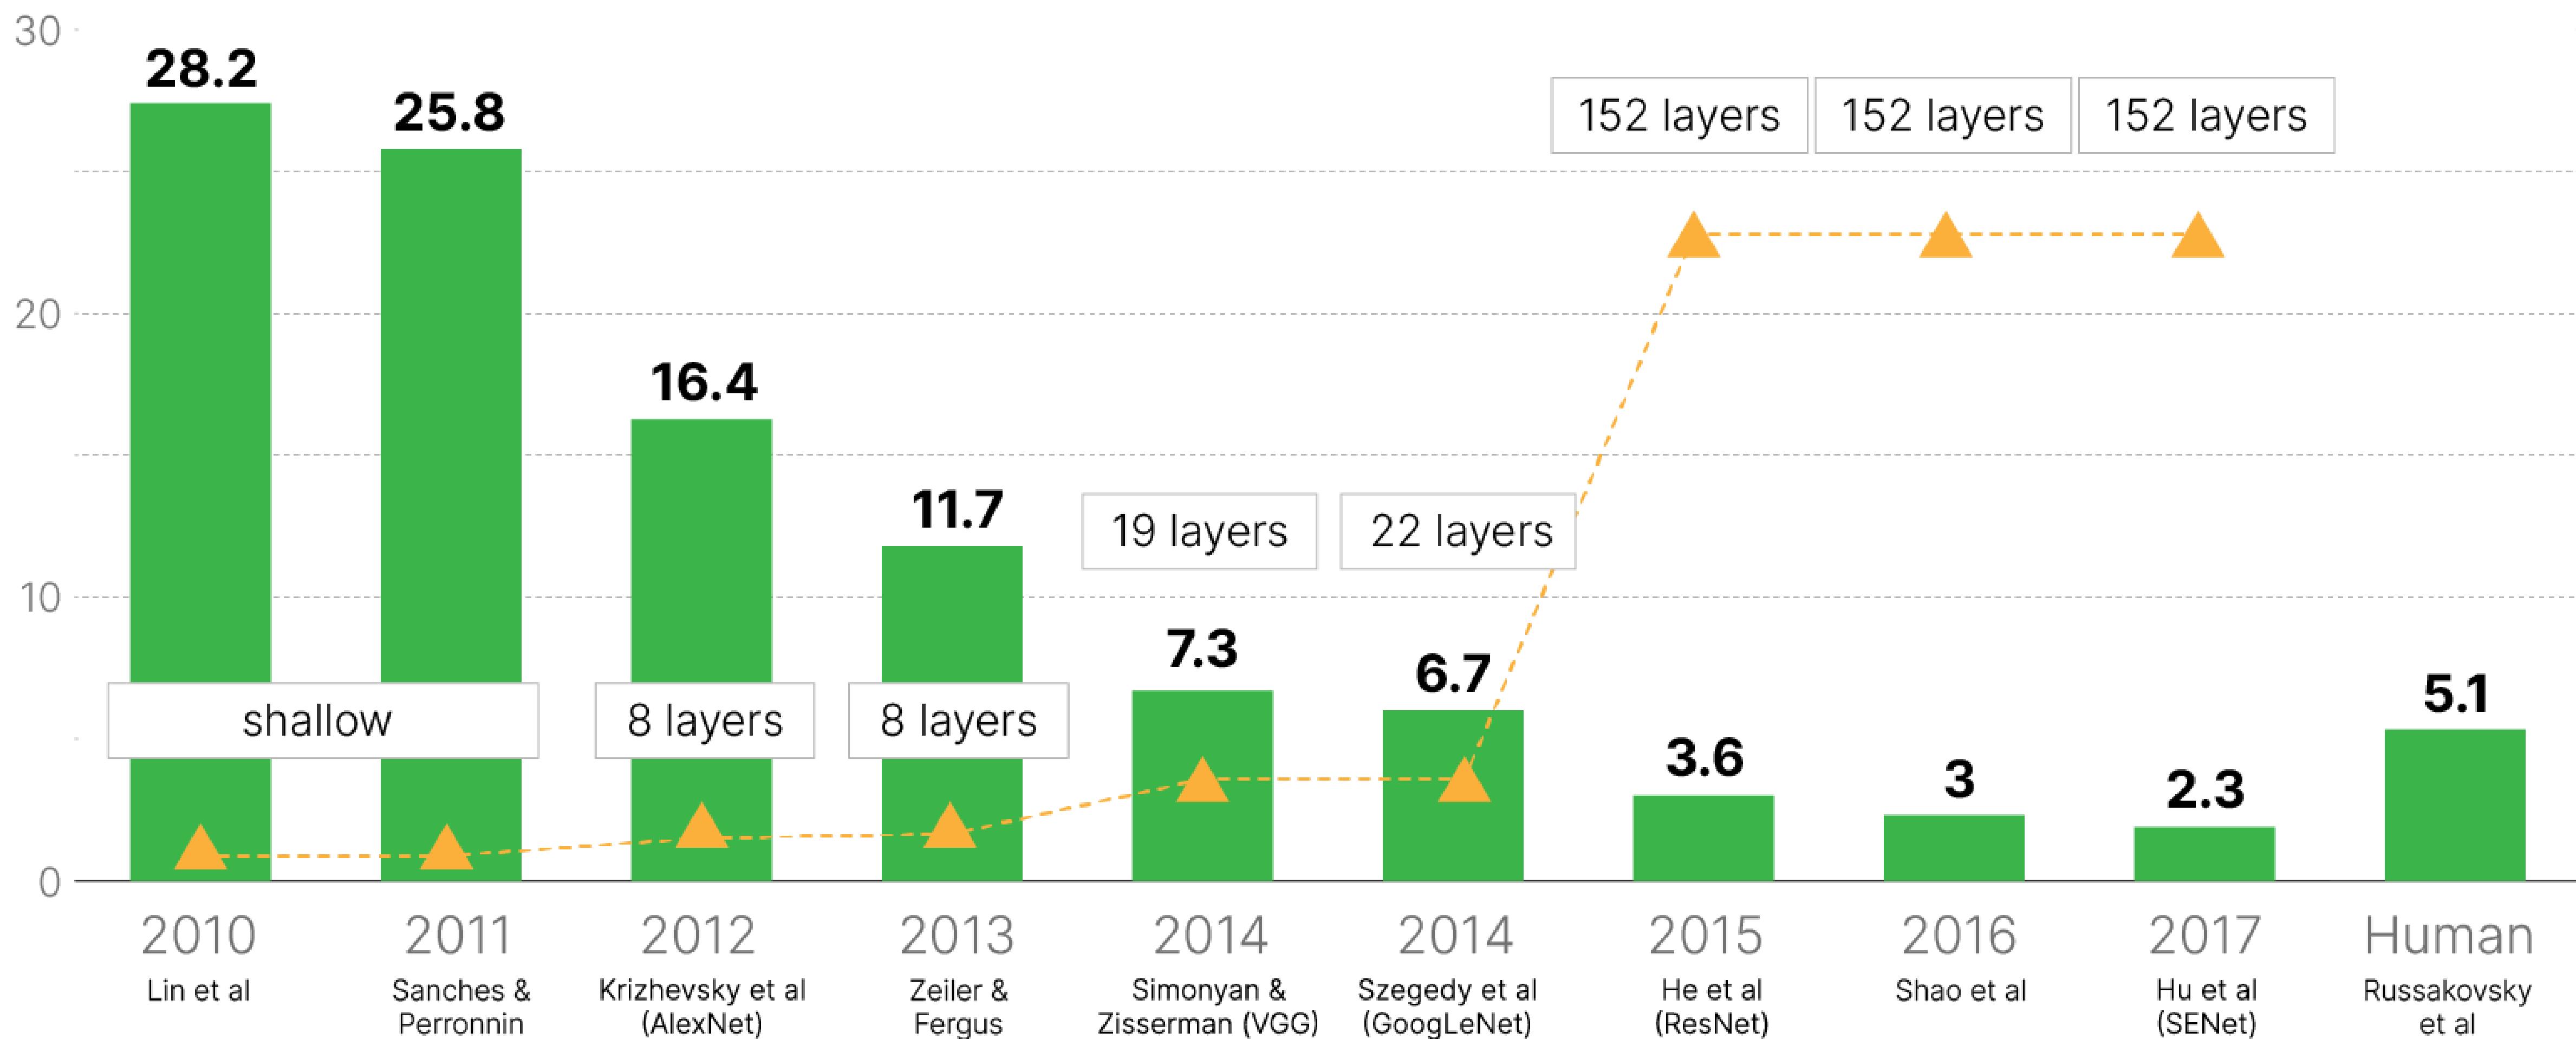
Conv Ops:

[1×1 conv, 64] $28 \times 28 \times 64 \times 1 \times 256$
[1×1 conv, 64] $28 \times 28 \times 64 \times 1 \times 256$
[1×1 conv, 128] $28 \times 28 \times 128 \times 1 \times 256$
[3×3 conv, 192] $28 \times 28 \times 192 \times 3 \times 64$
[5×5 conv, 96] $28 \times 28 \times 96 \times 5 \times 64$
[1×1 conv, 64] $28 \times 28 \times 64 \times 1 \times 256$
Total: 358M ops

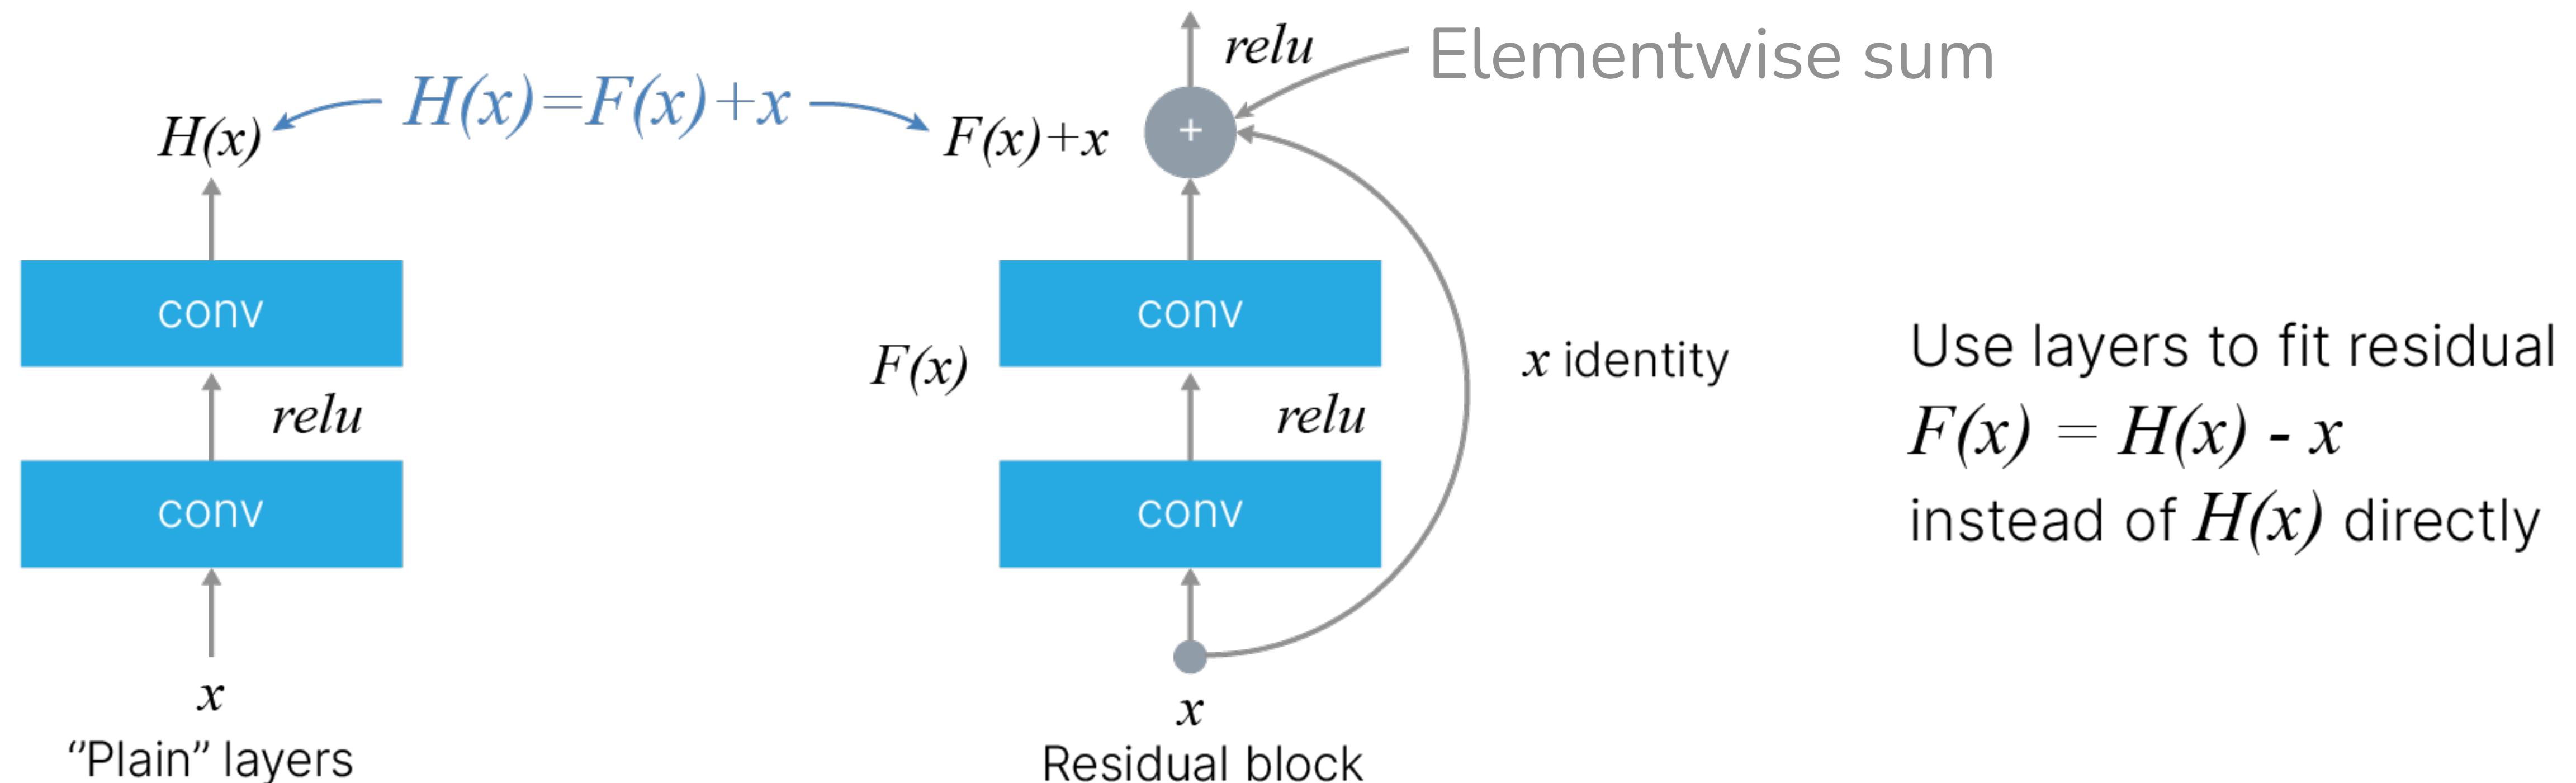
GoogLeNet (2014) | Vanishing gradients



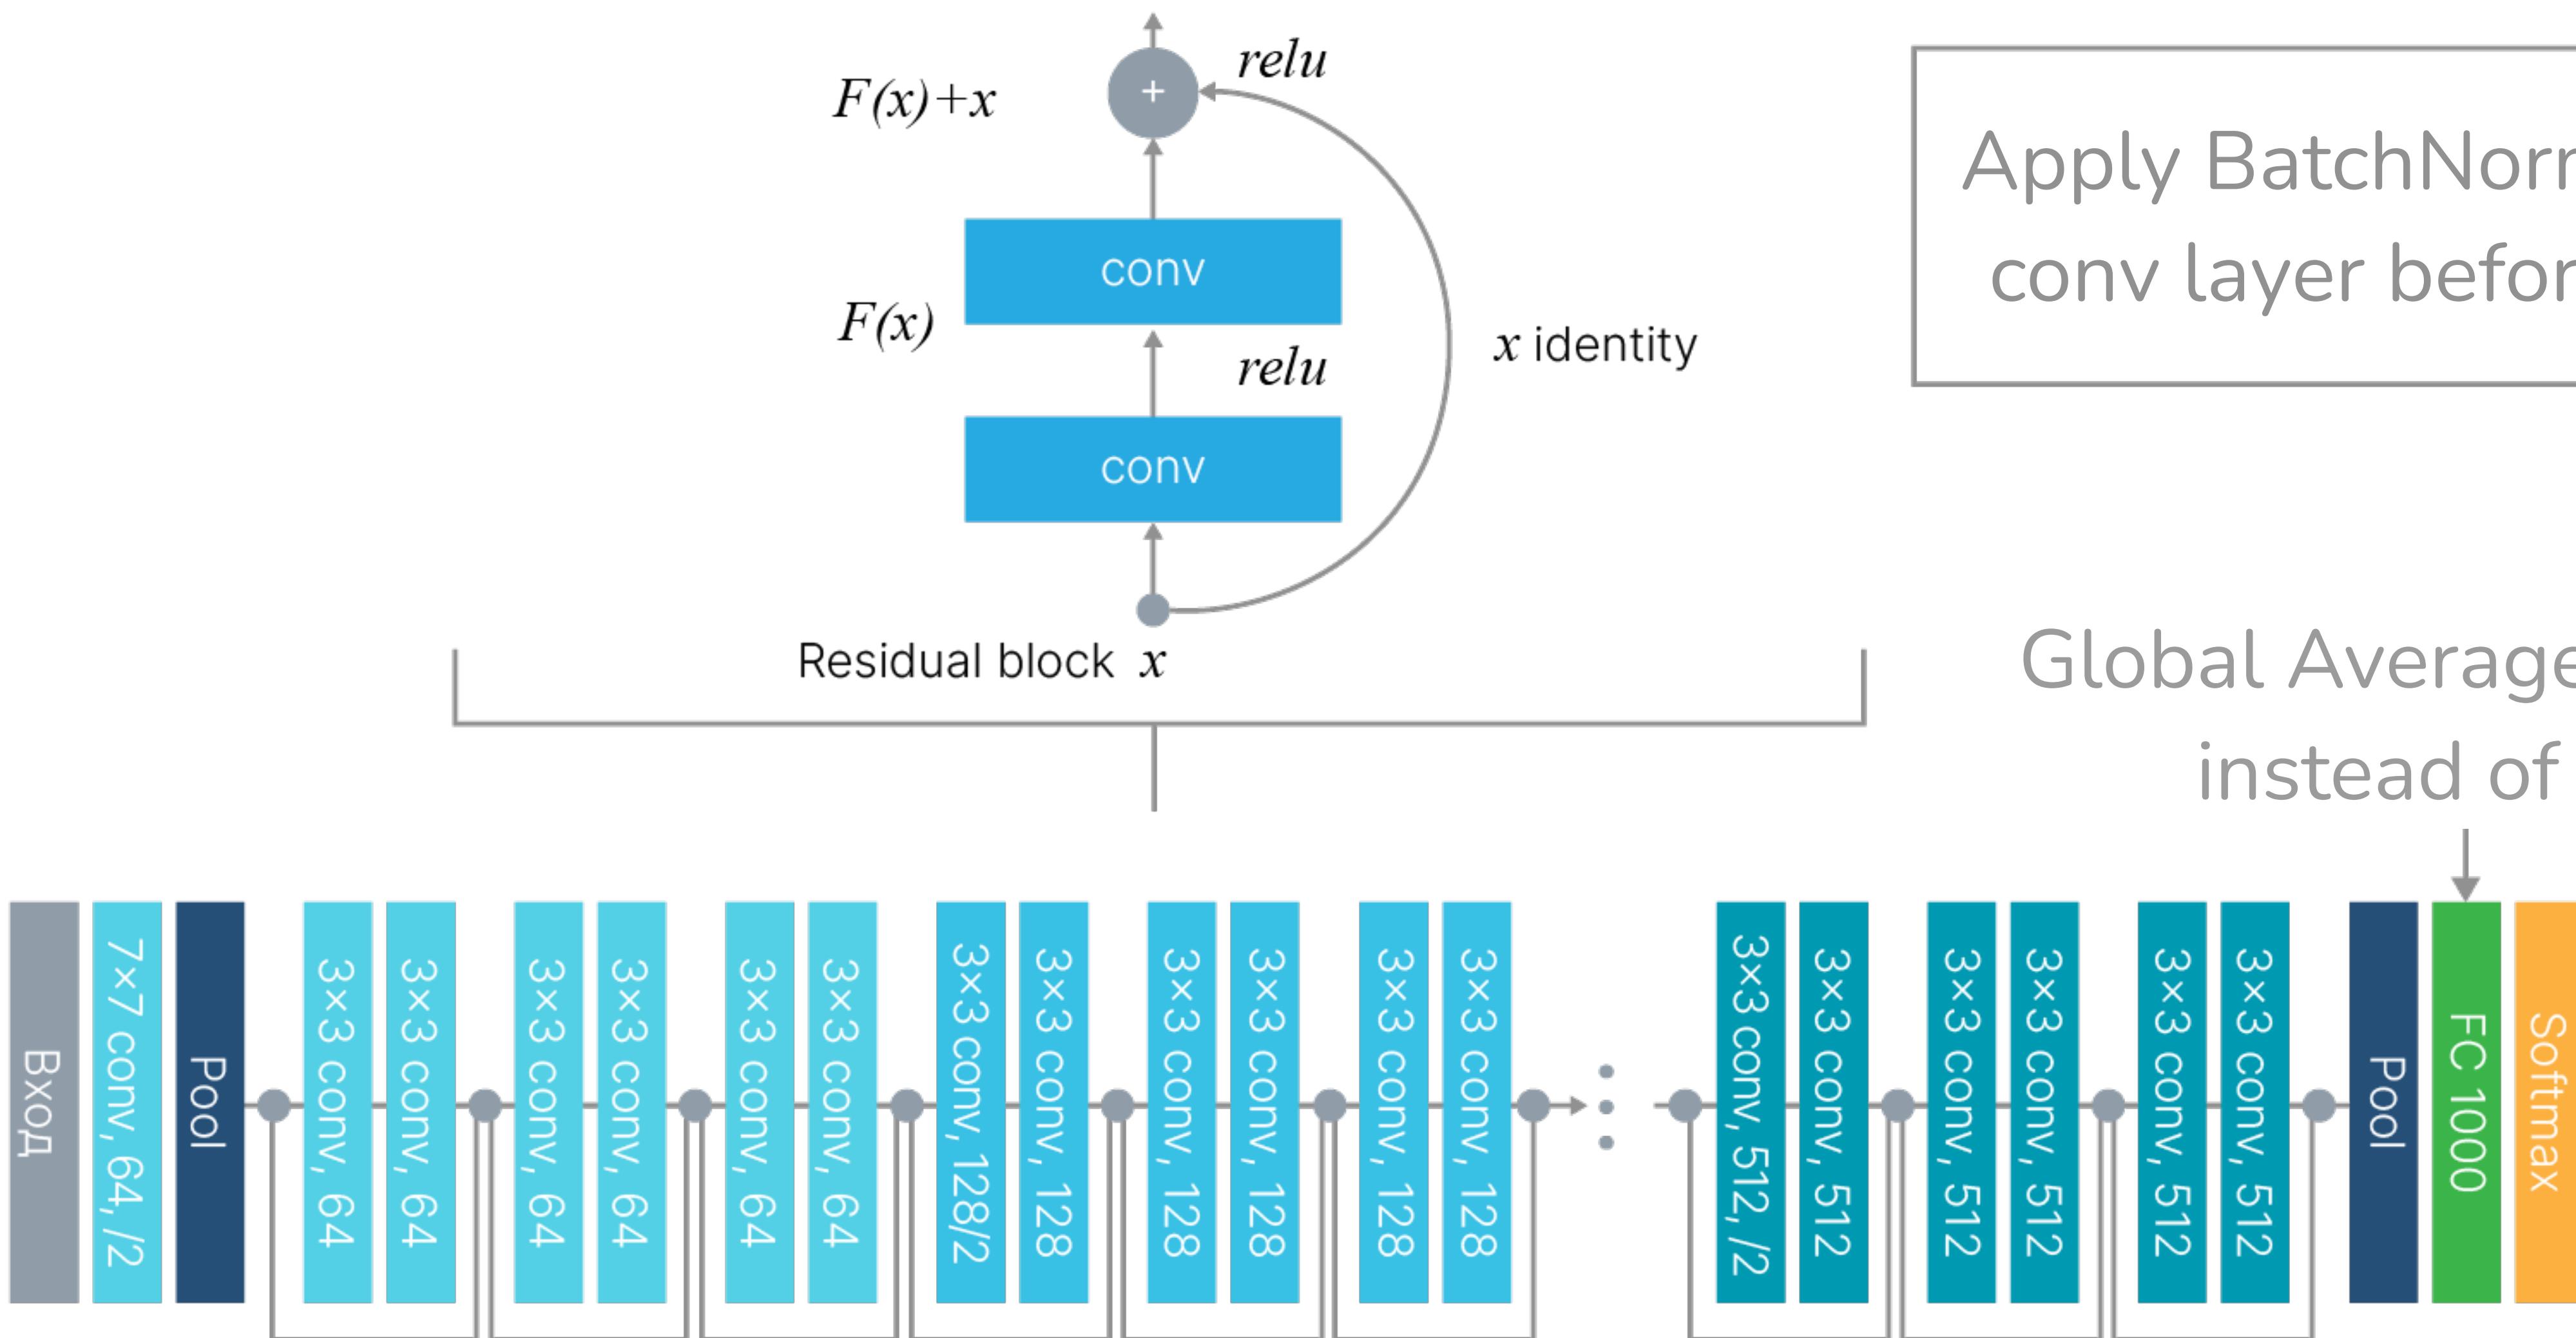
ImageNet Large Scale Visual Recognition Challenge



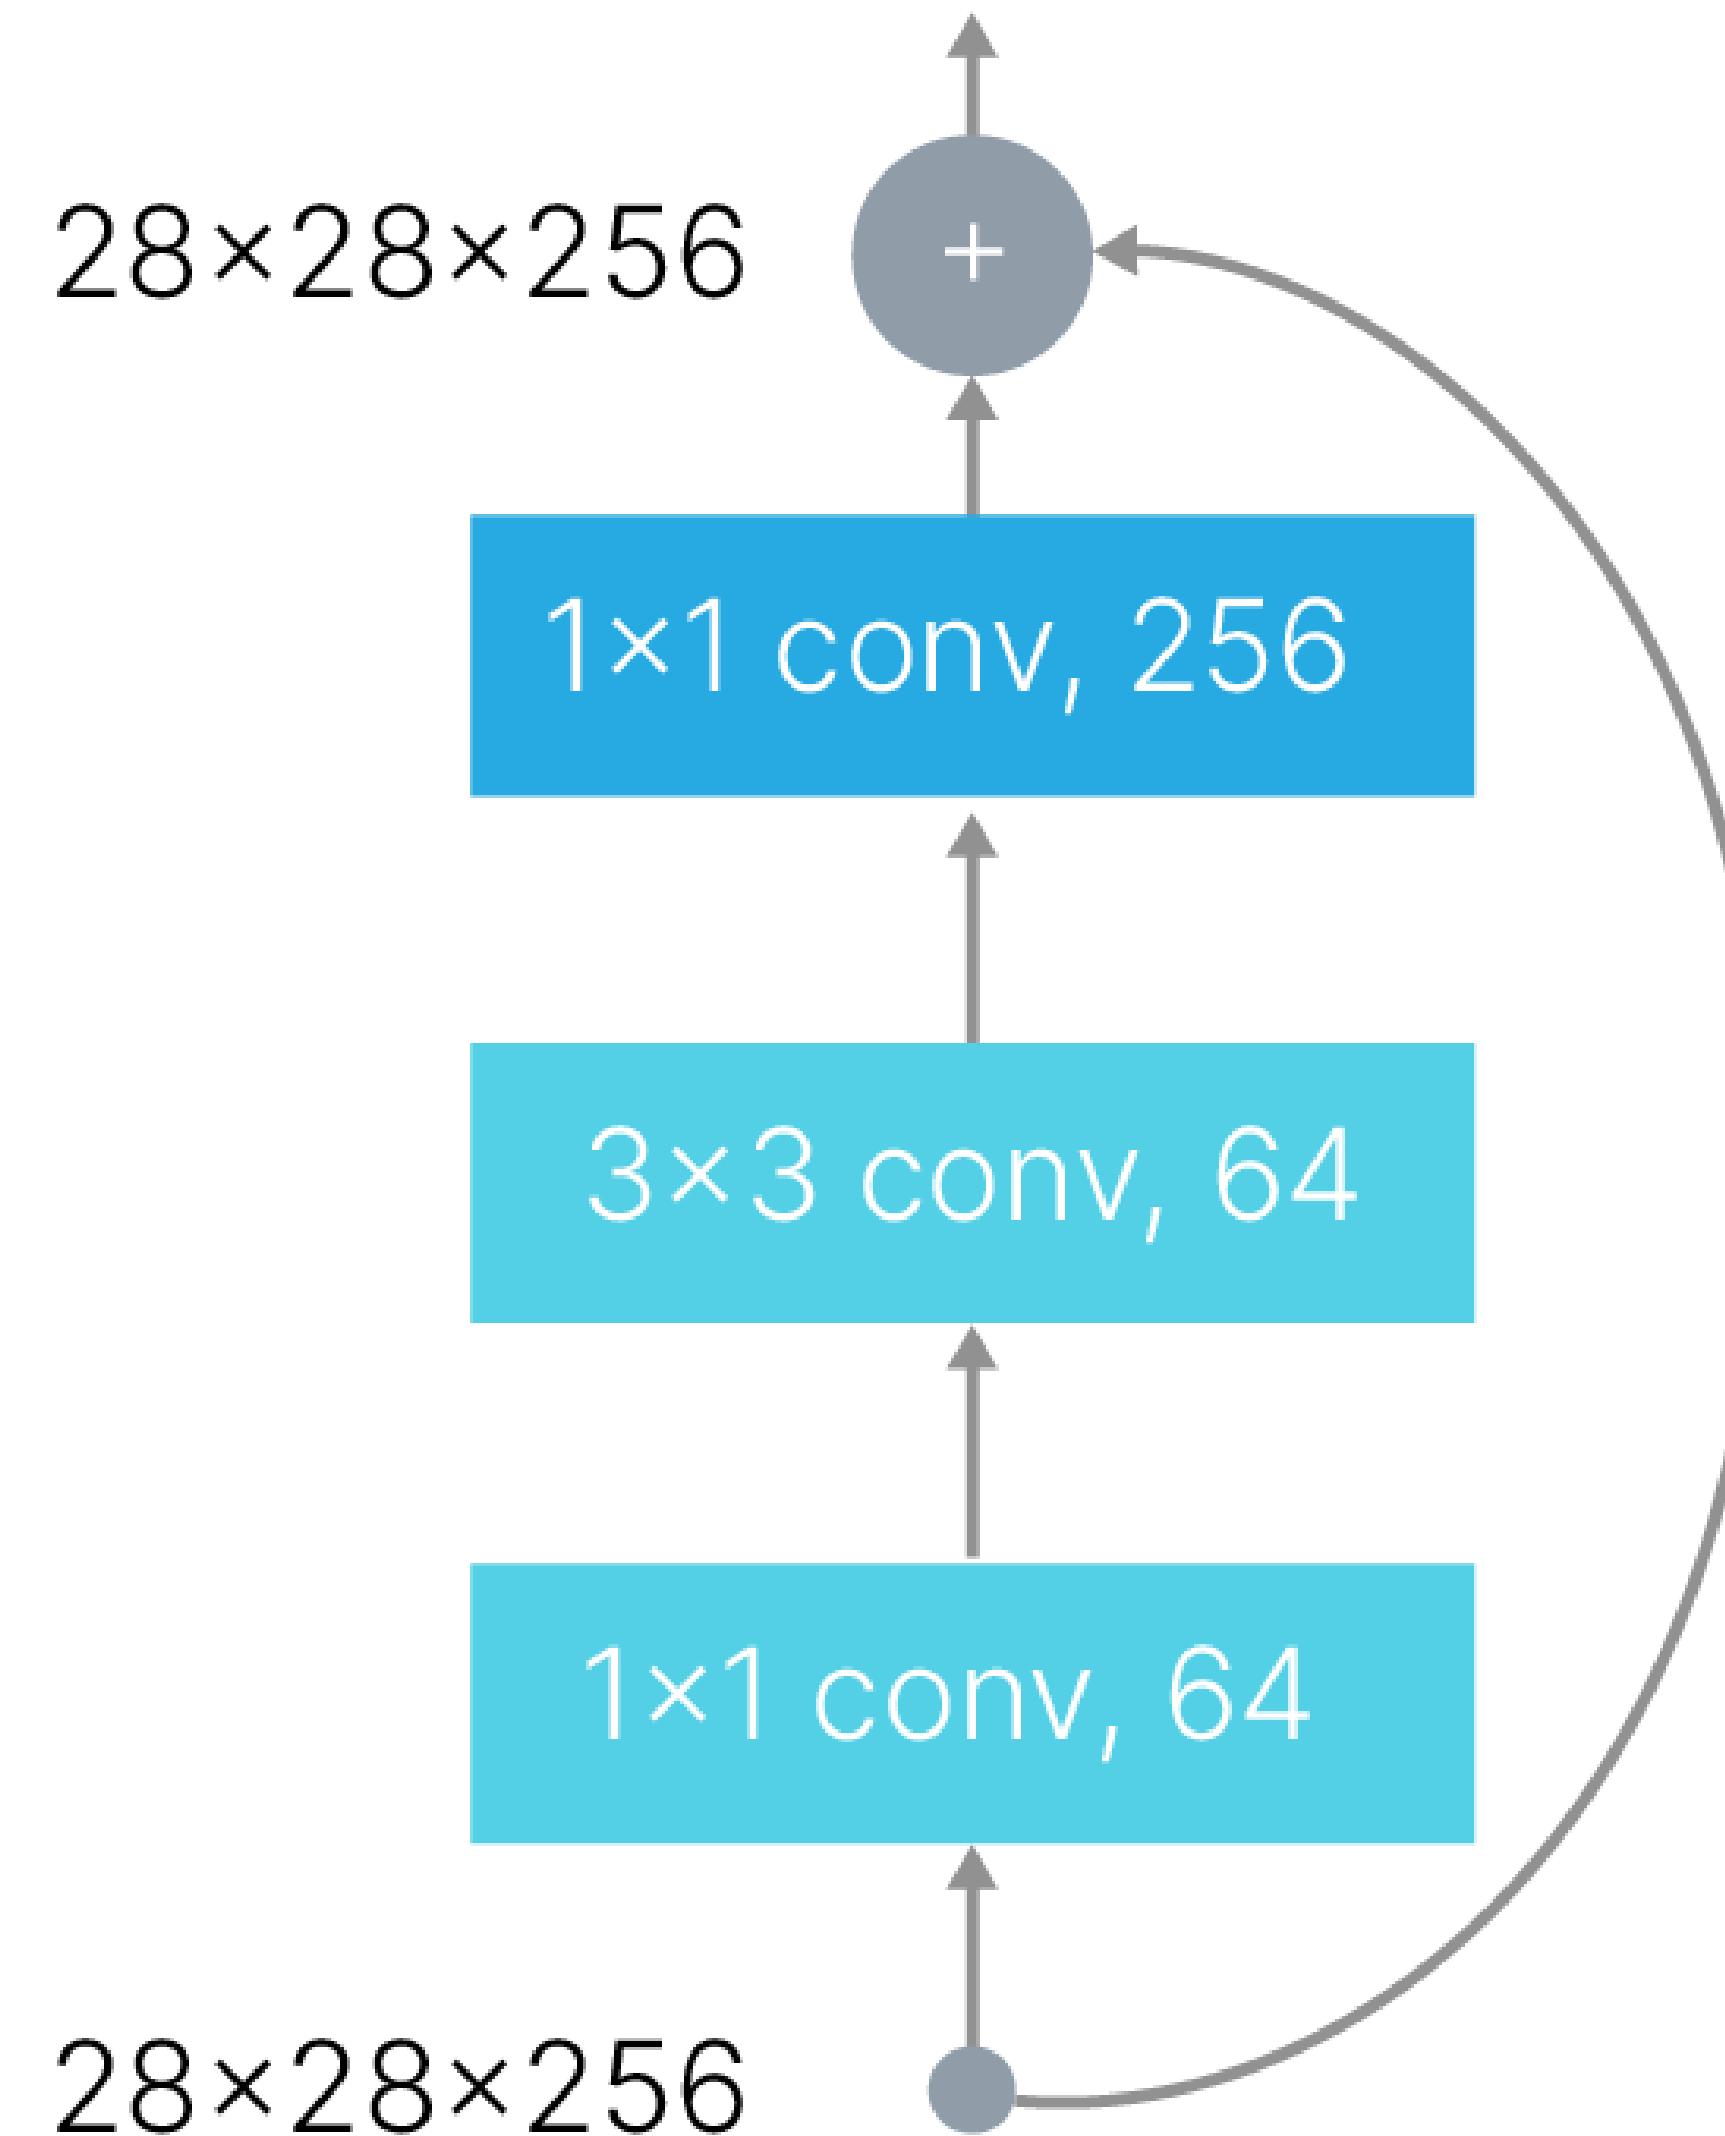
ResNet (2015) | Residual Connection



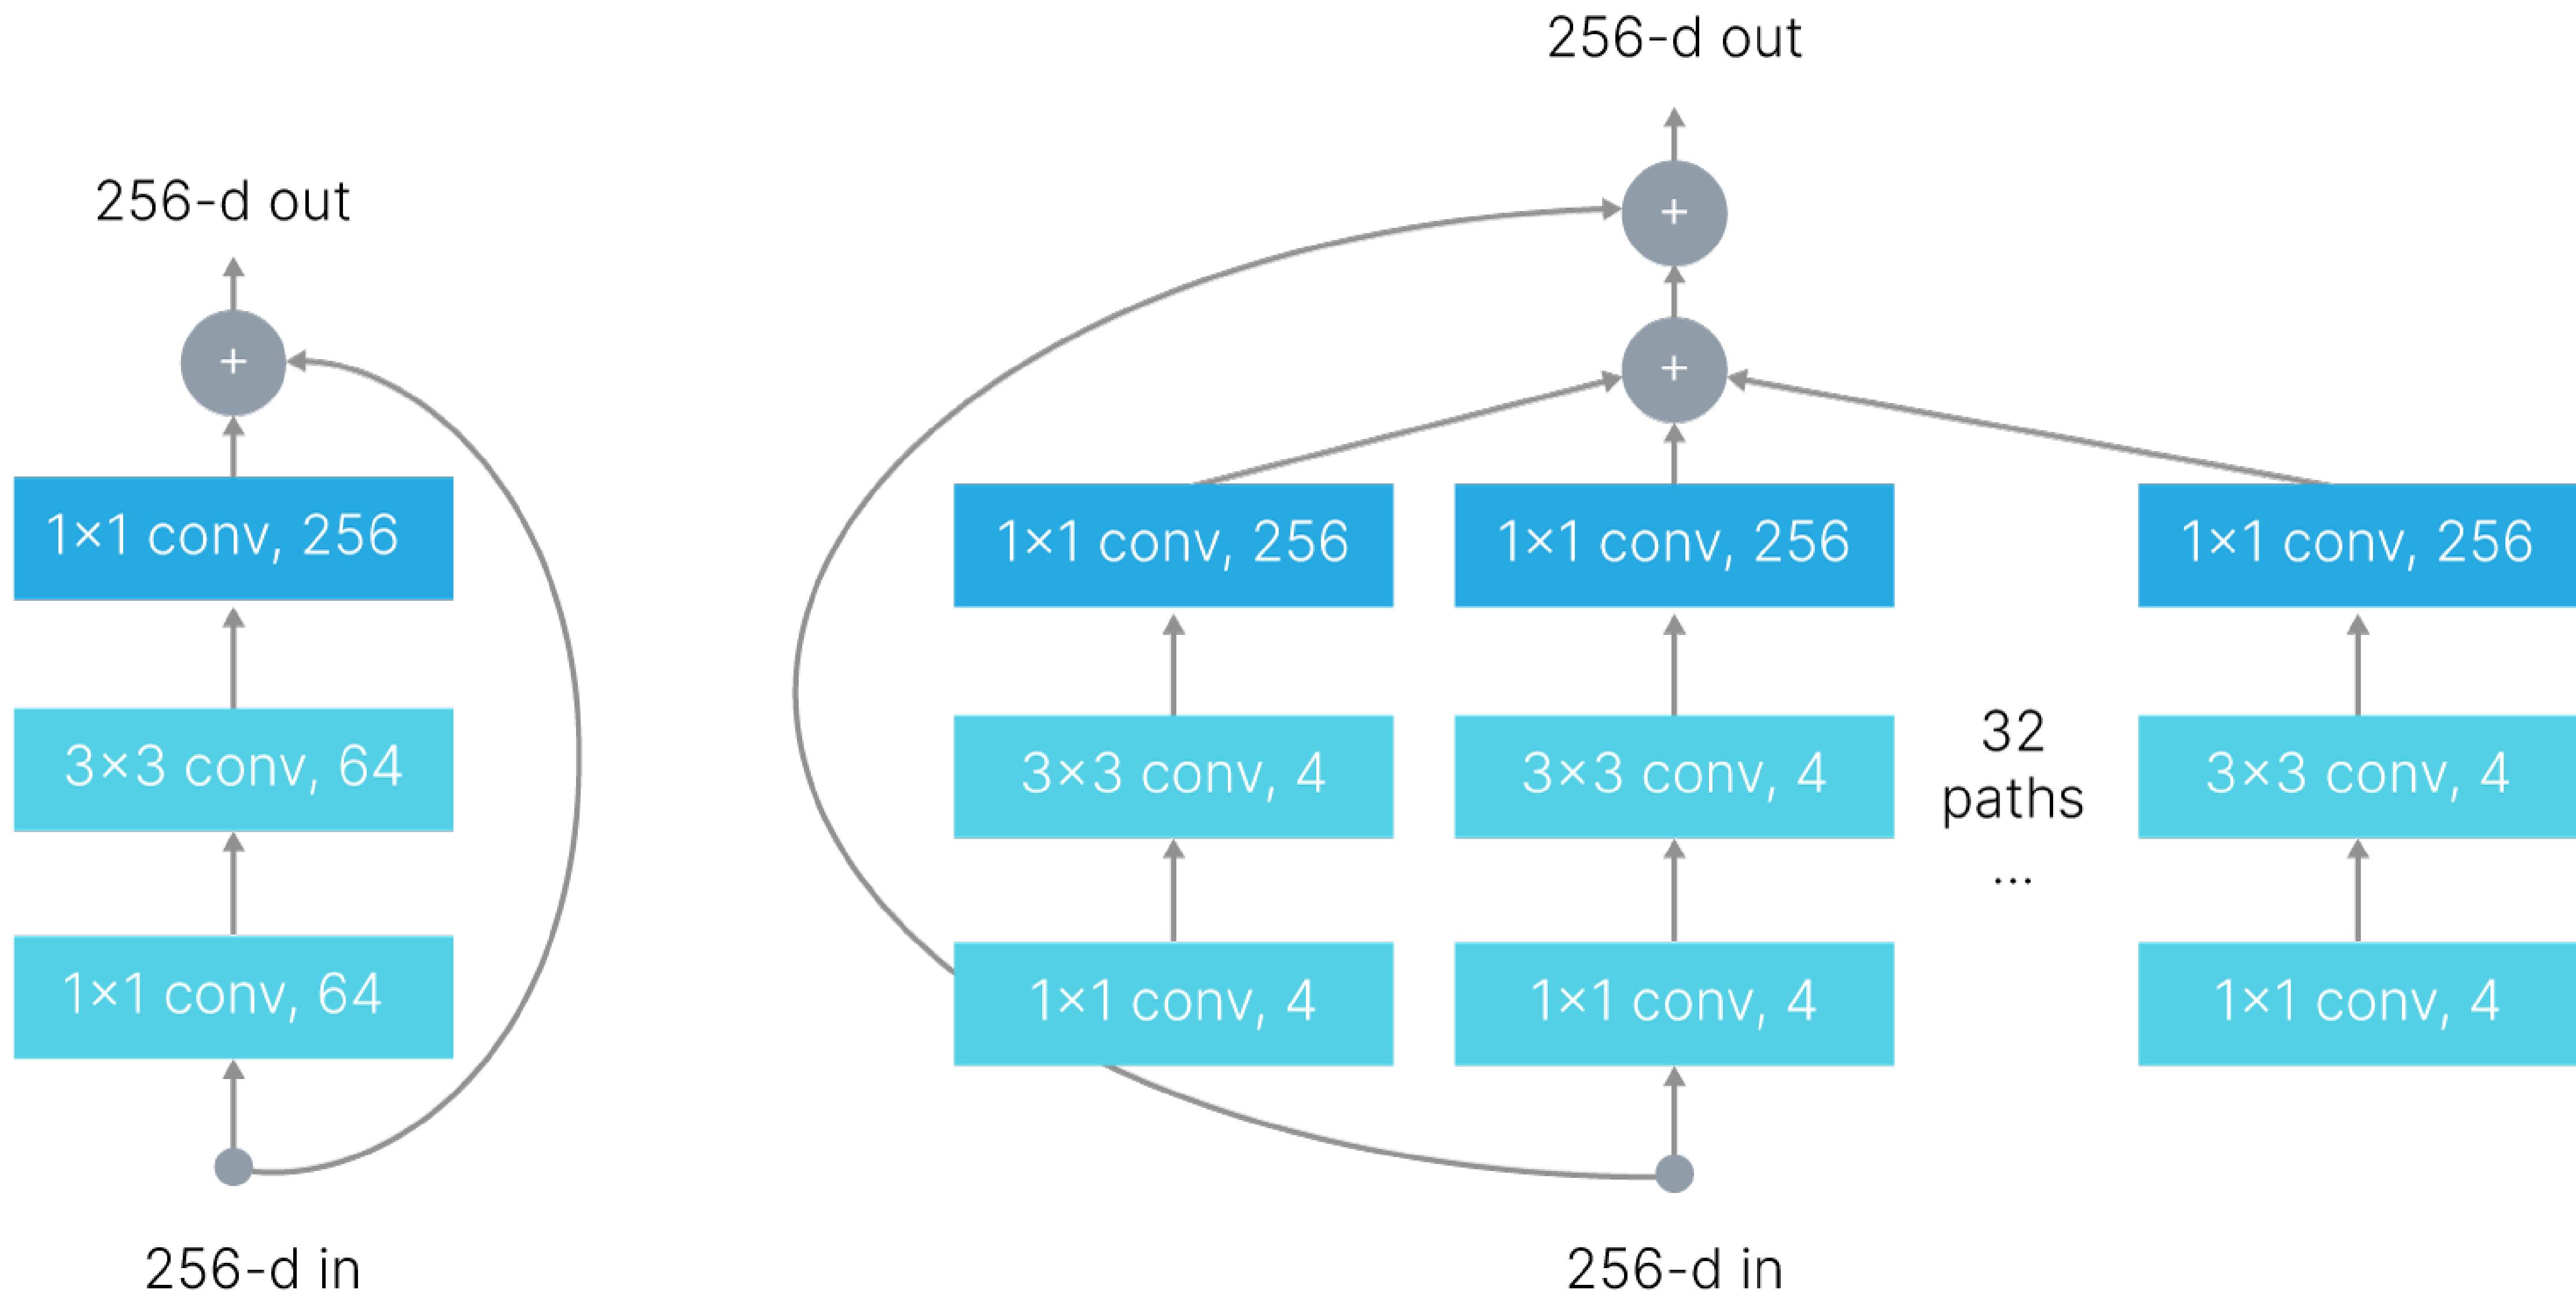
ResNet (2015) | Architecture



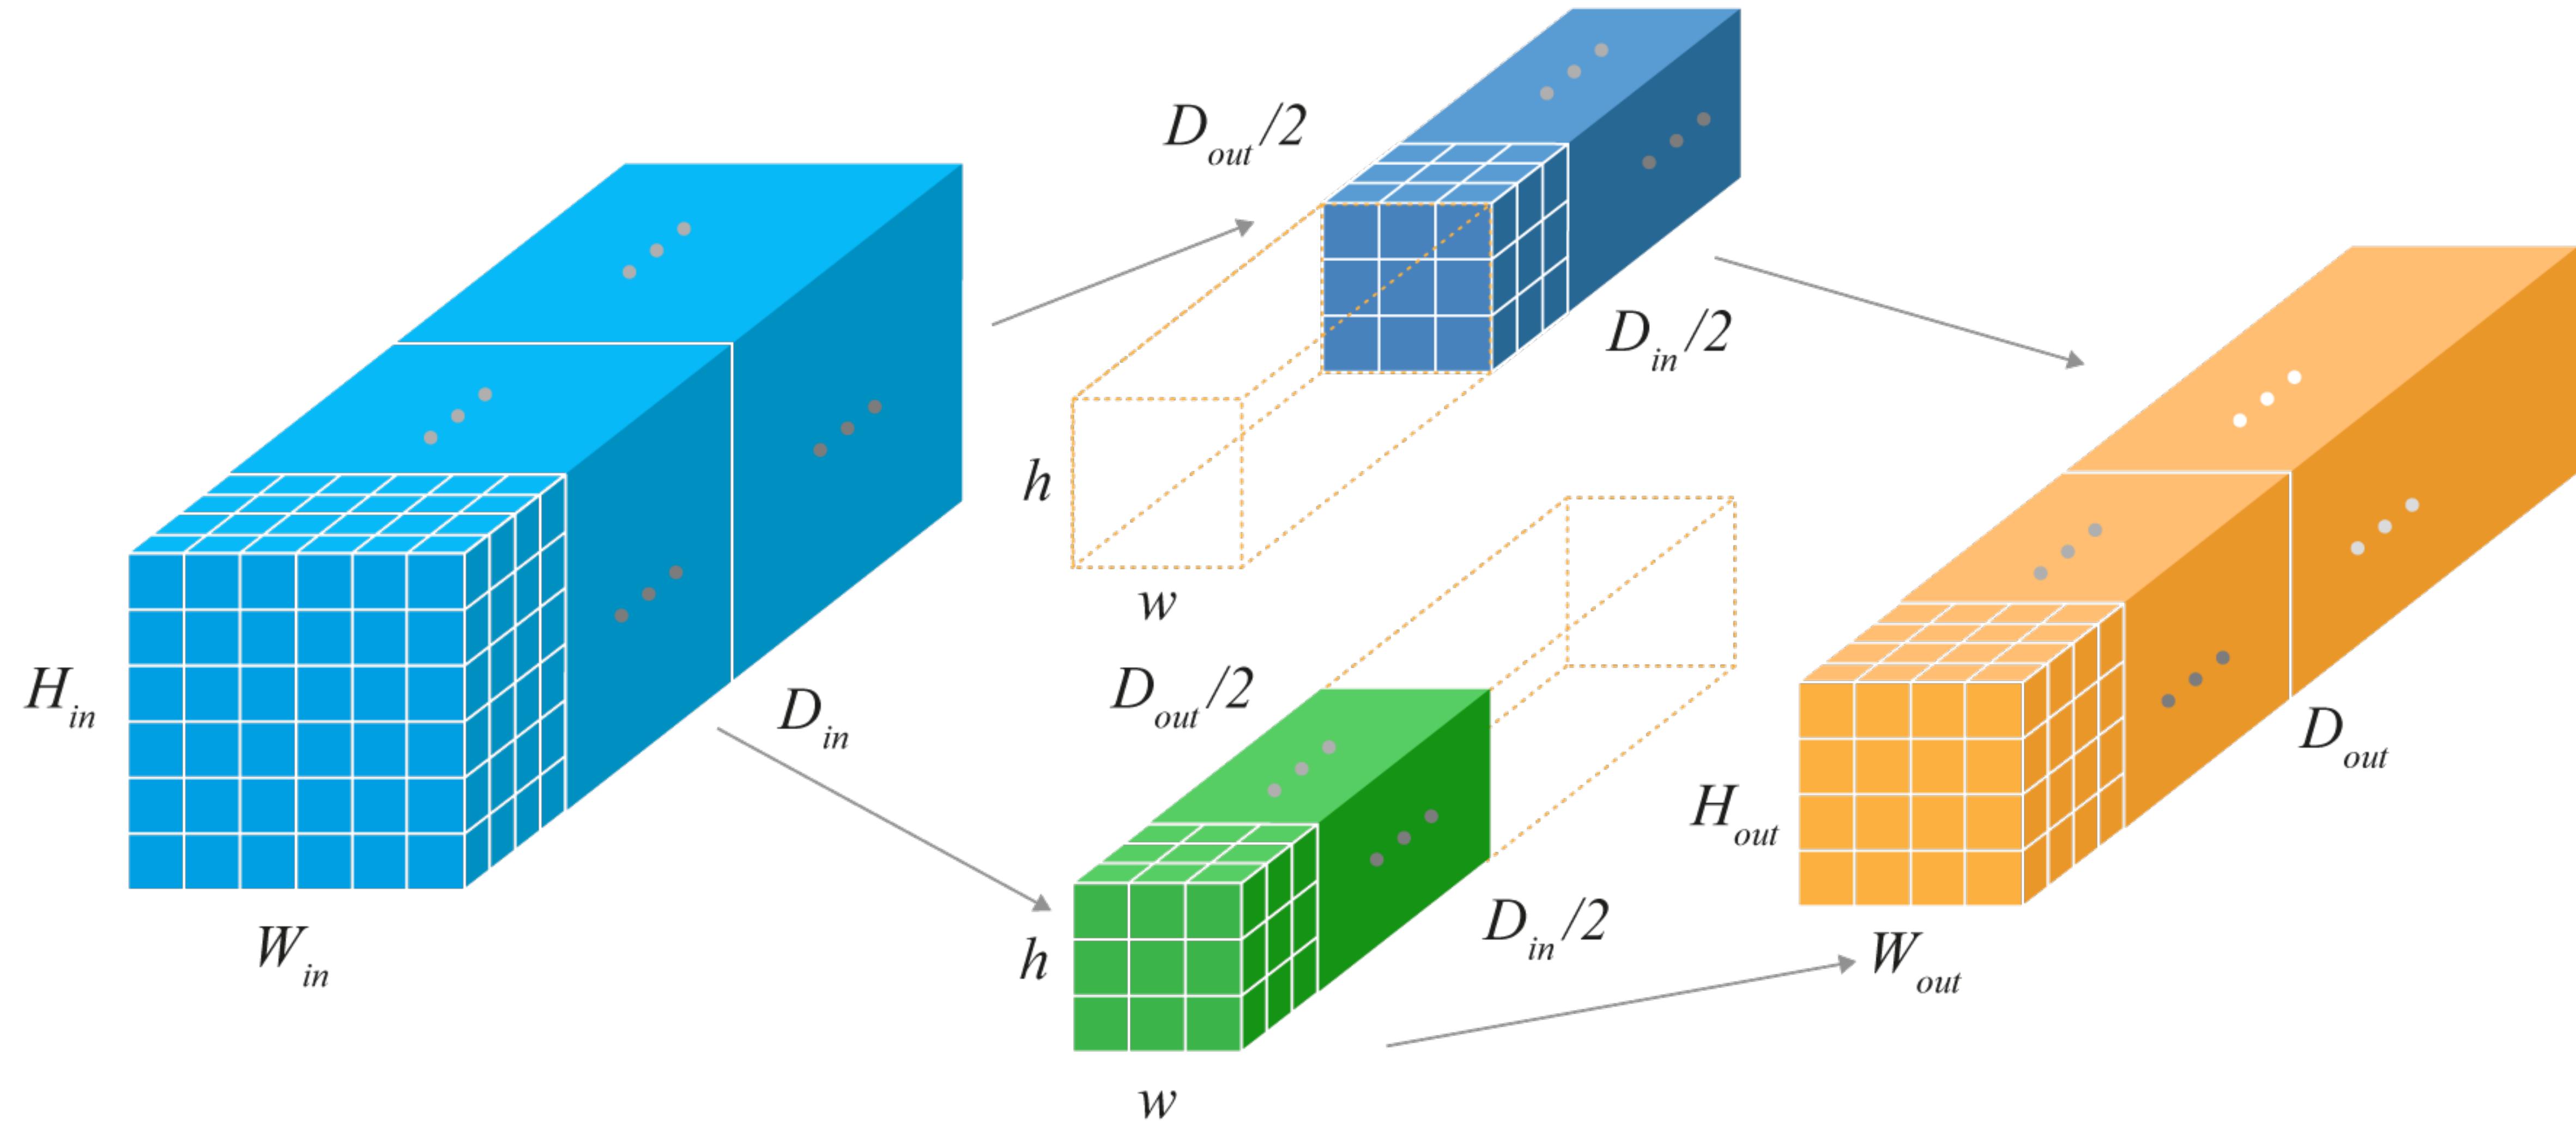
ResNet (2015) | Bottleneck layer



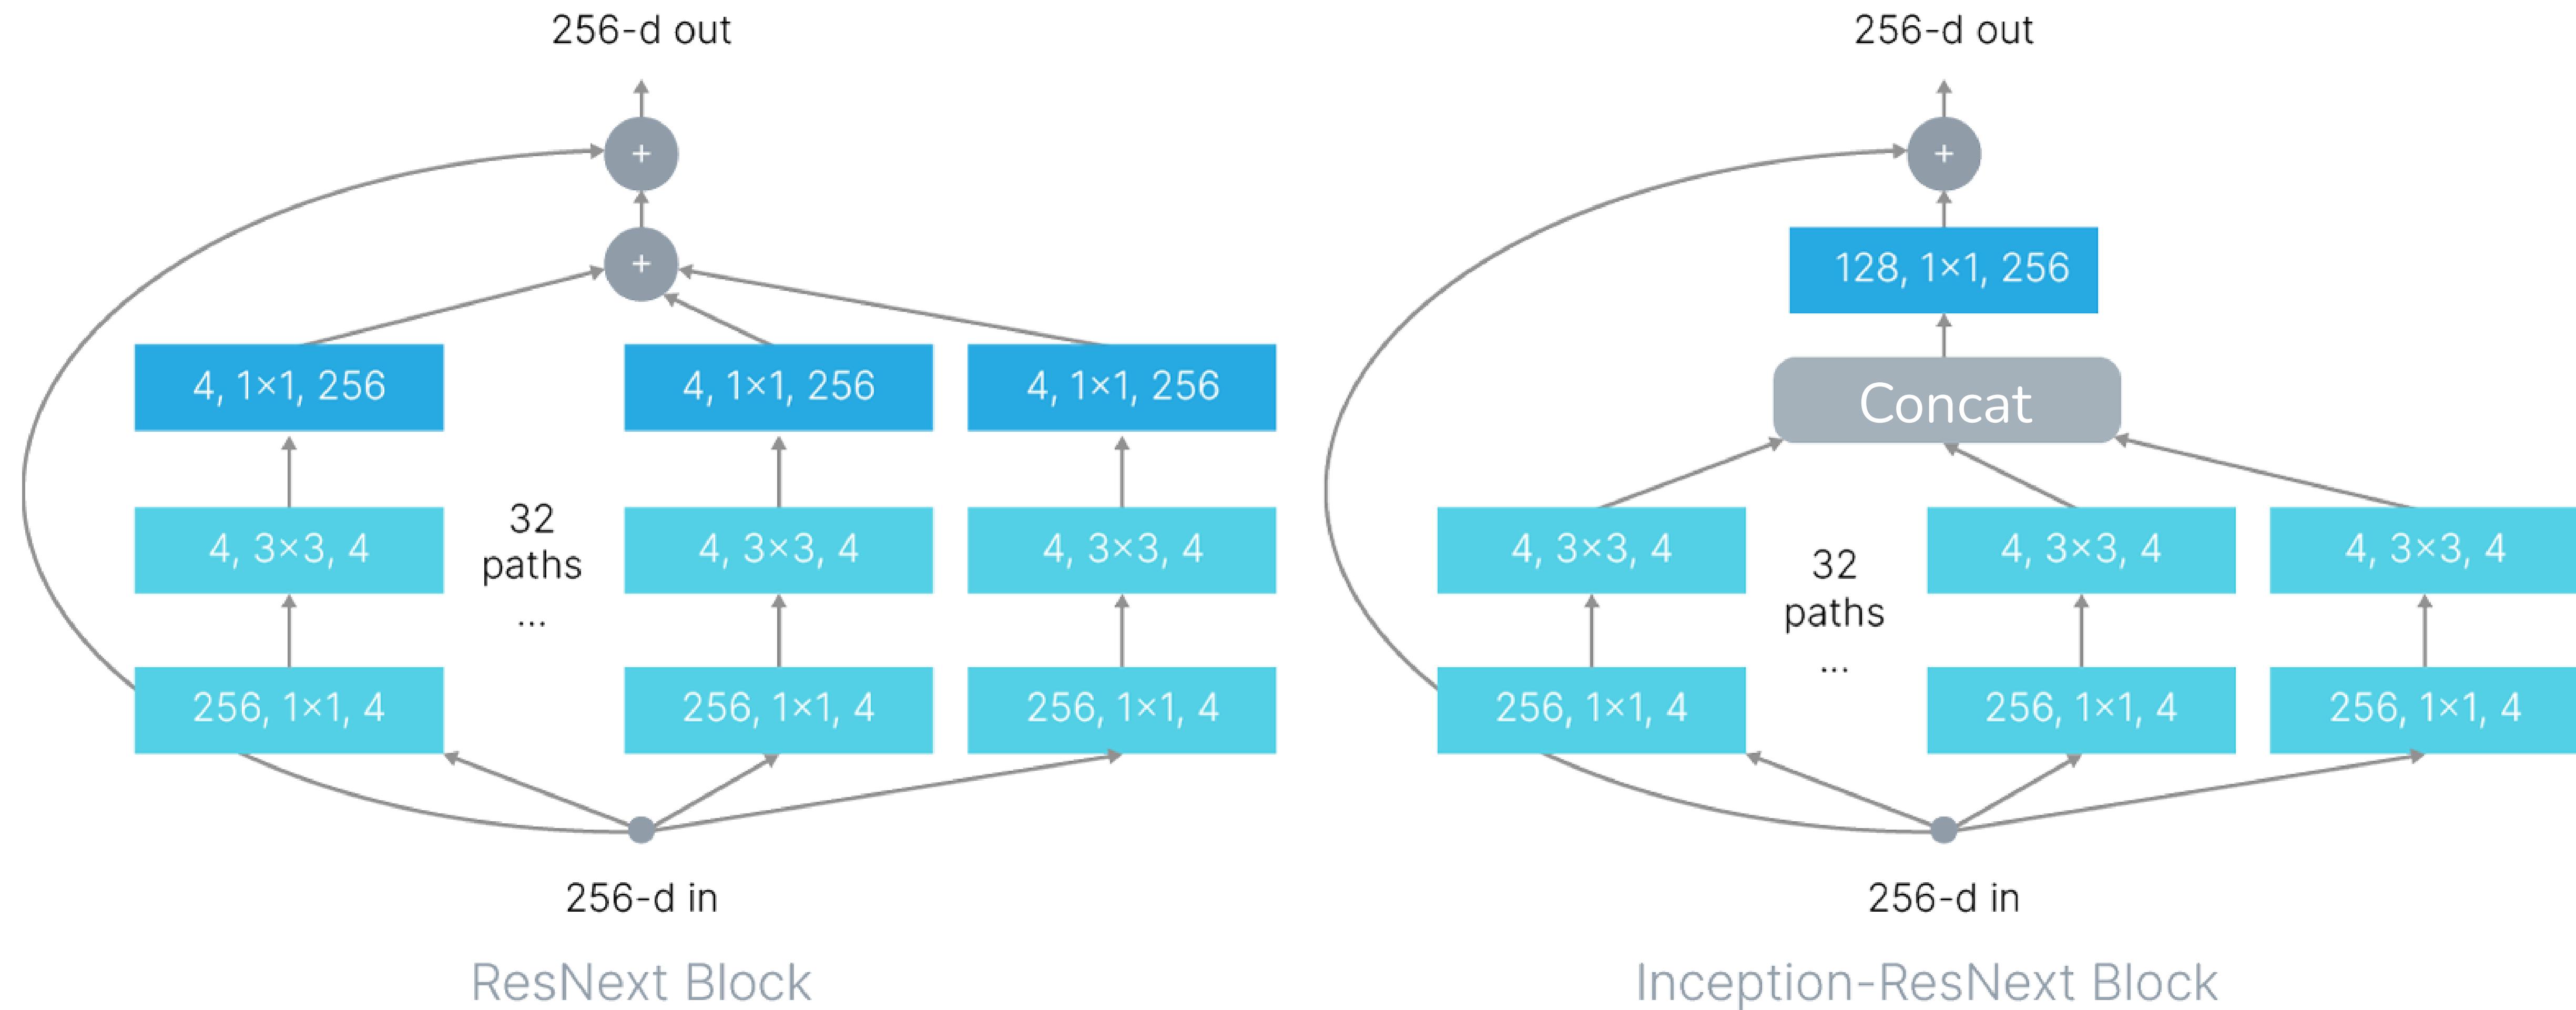
ResNeXt (2015) | Basic Block



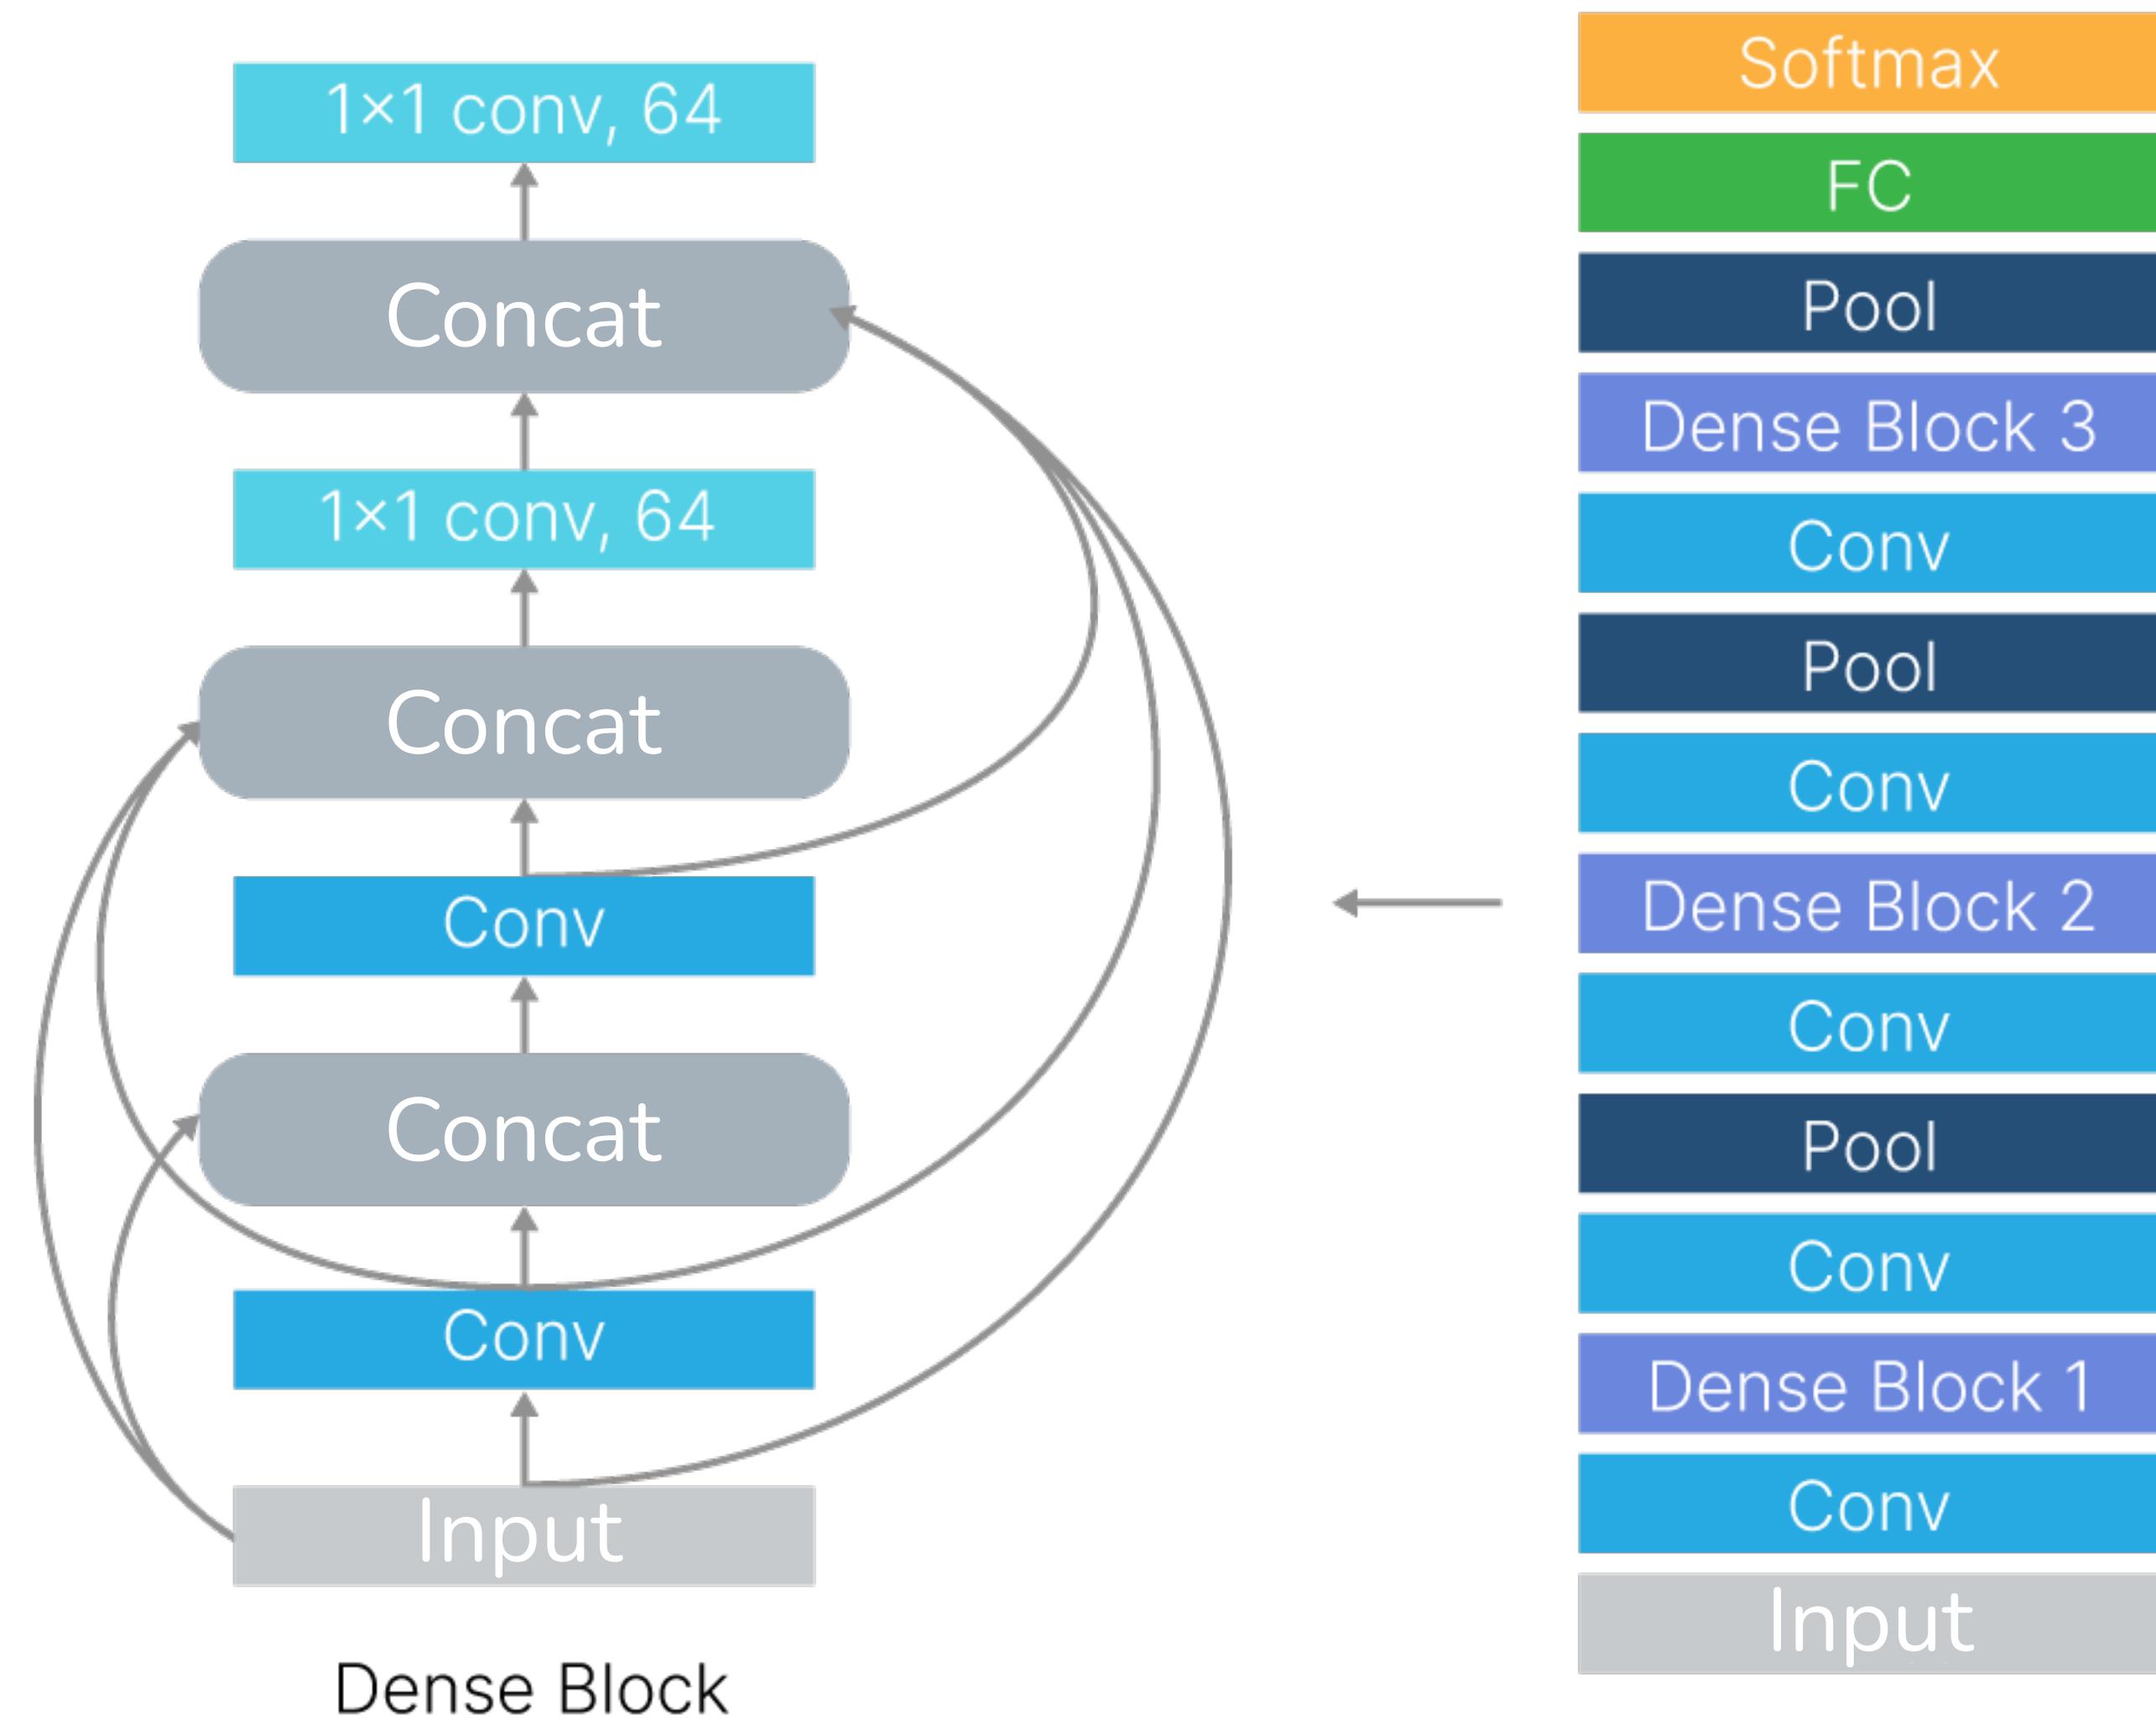
ResNeXt (2015) | Grouped Convolutions



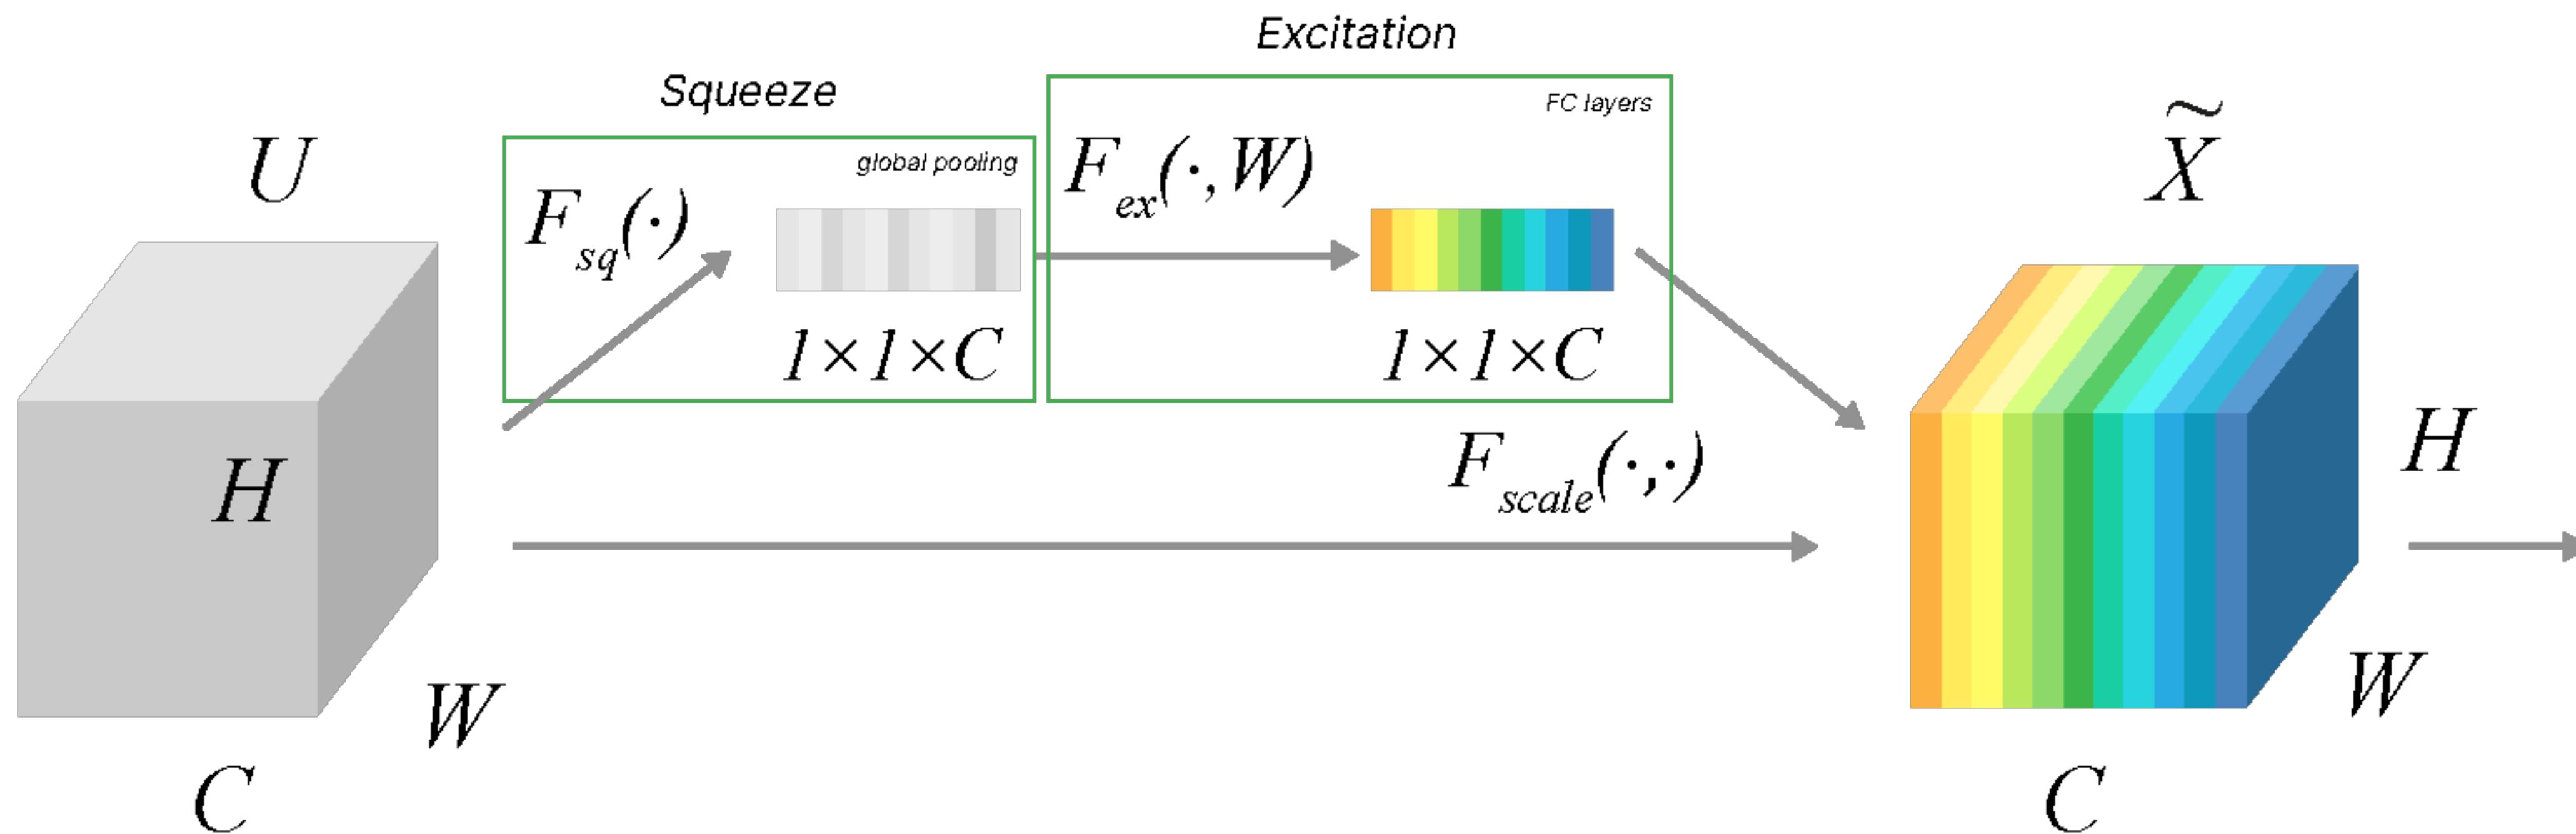
Inception-ResNeXt



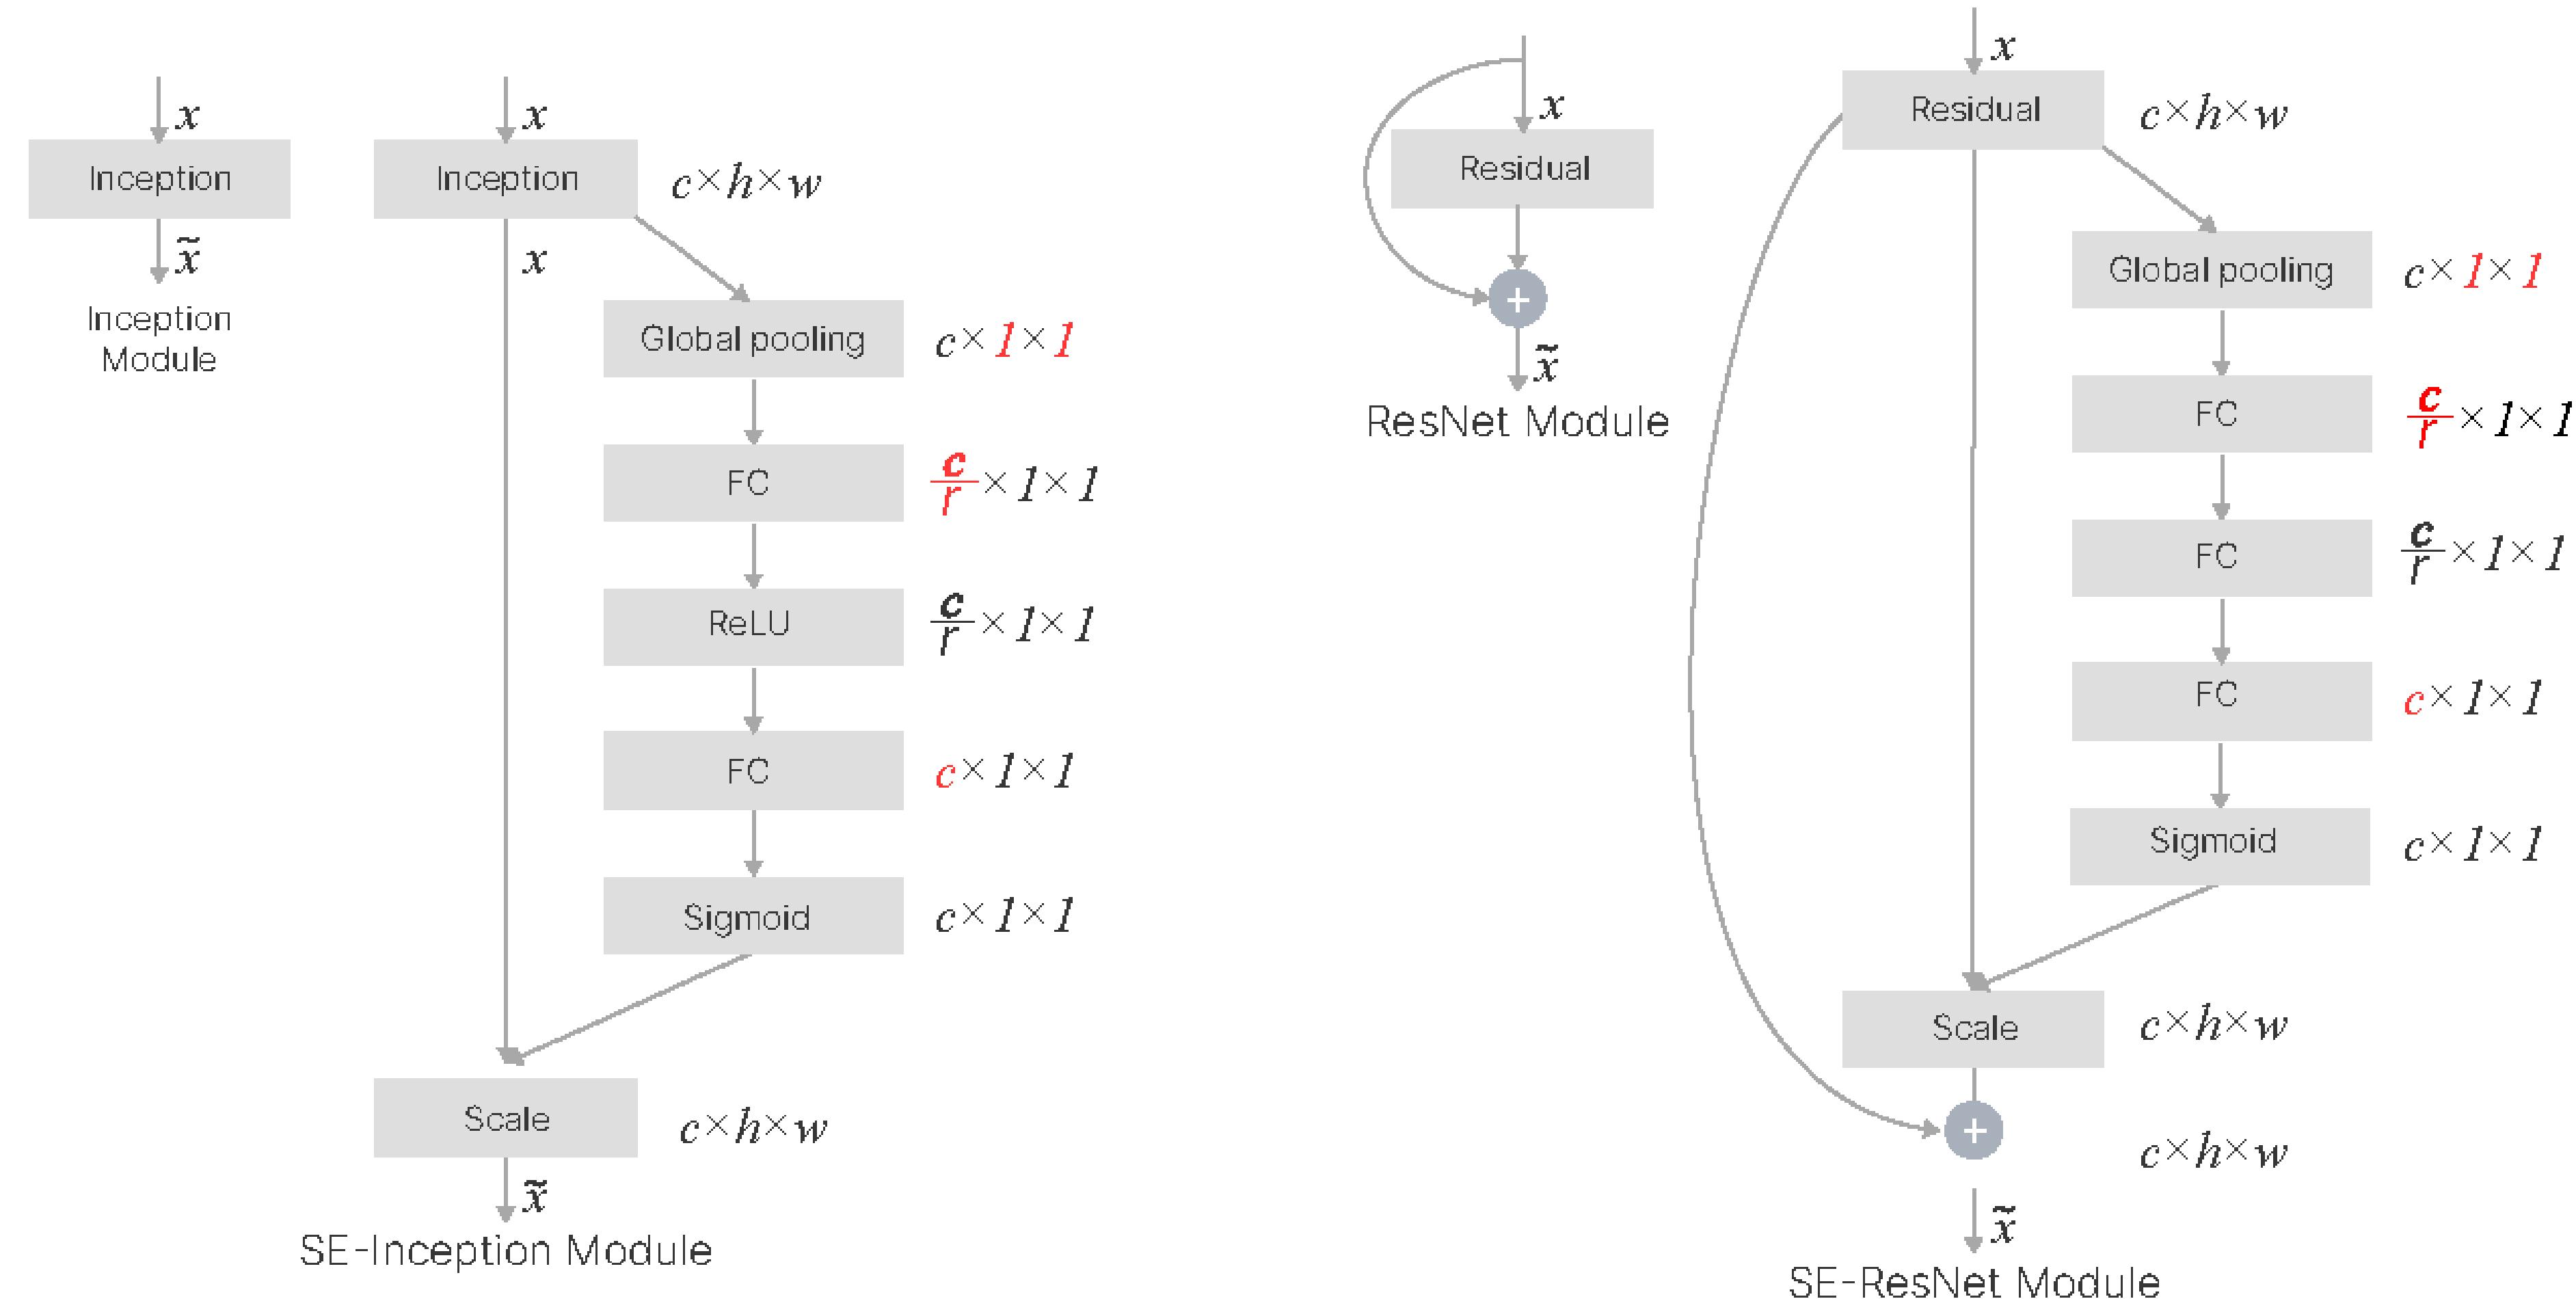
DenseNet (2016) | Architecture



SENet (2017) | Building block

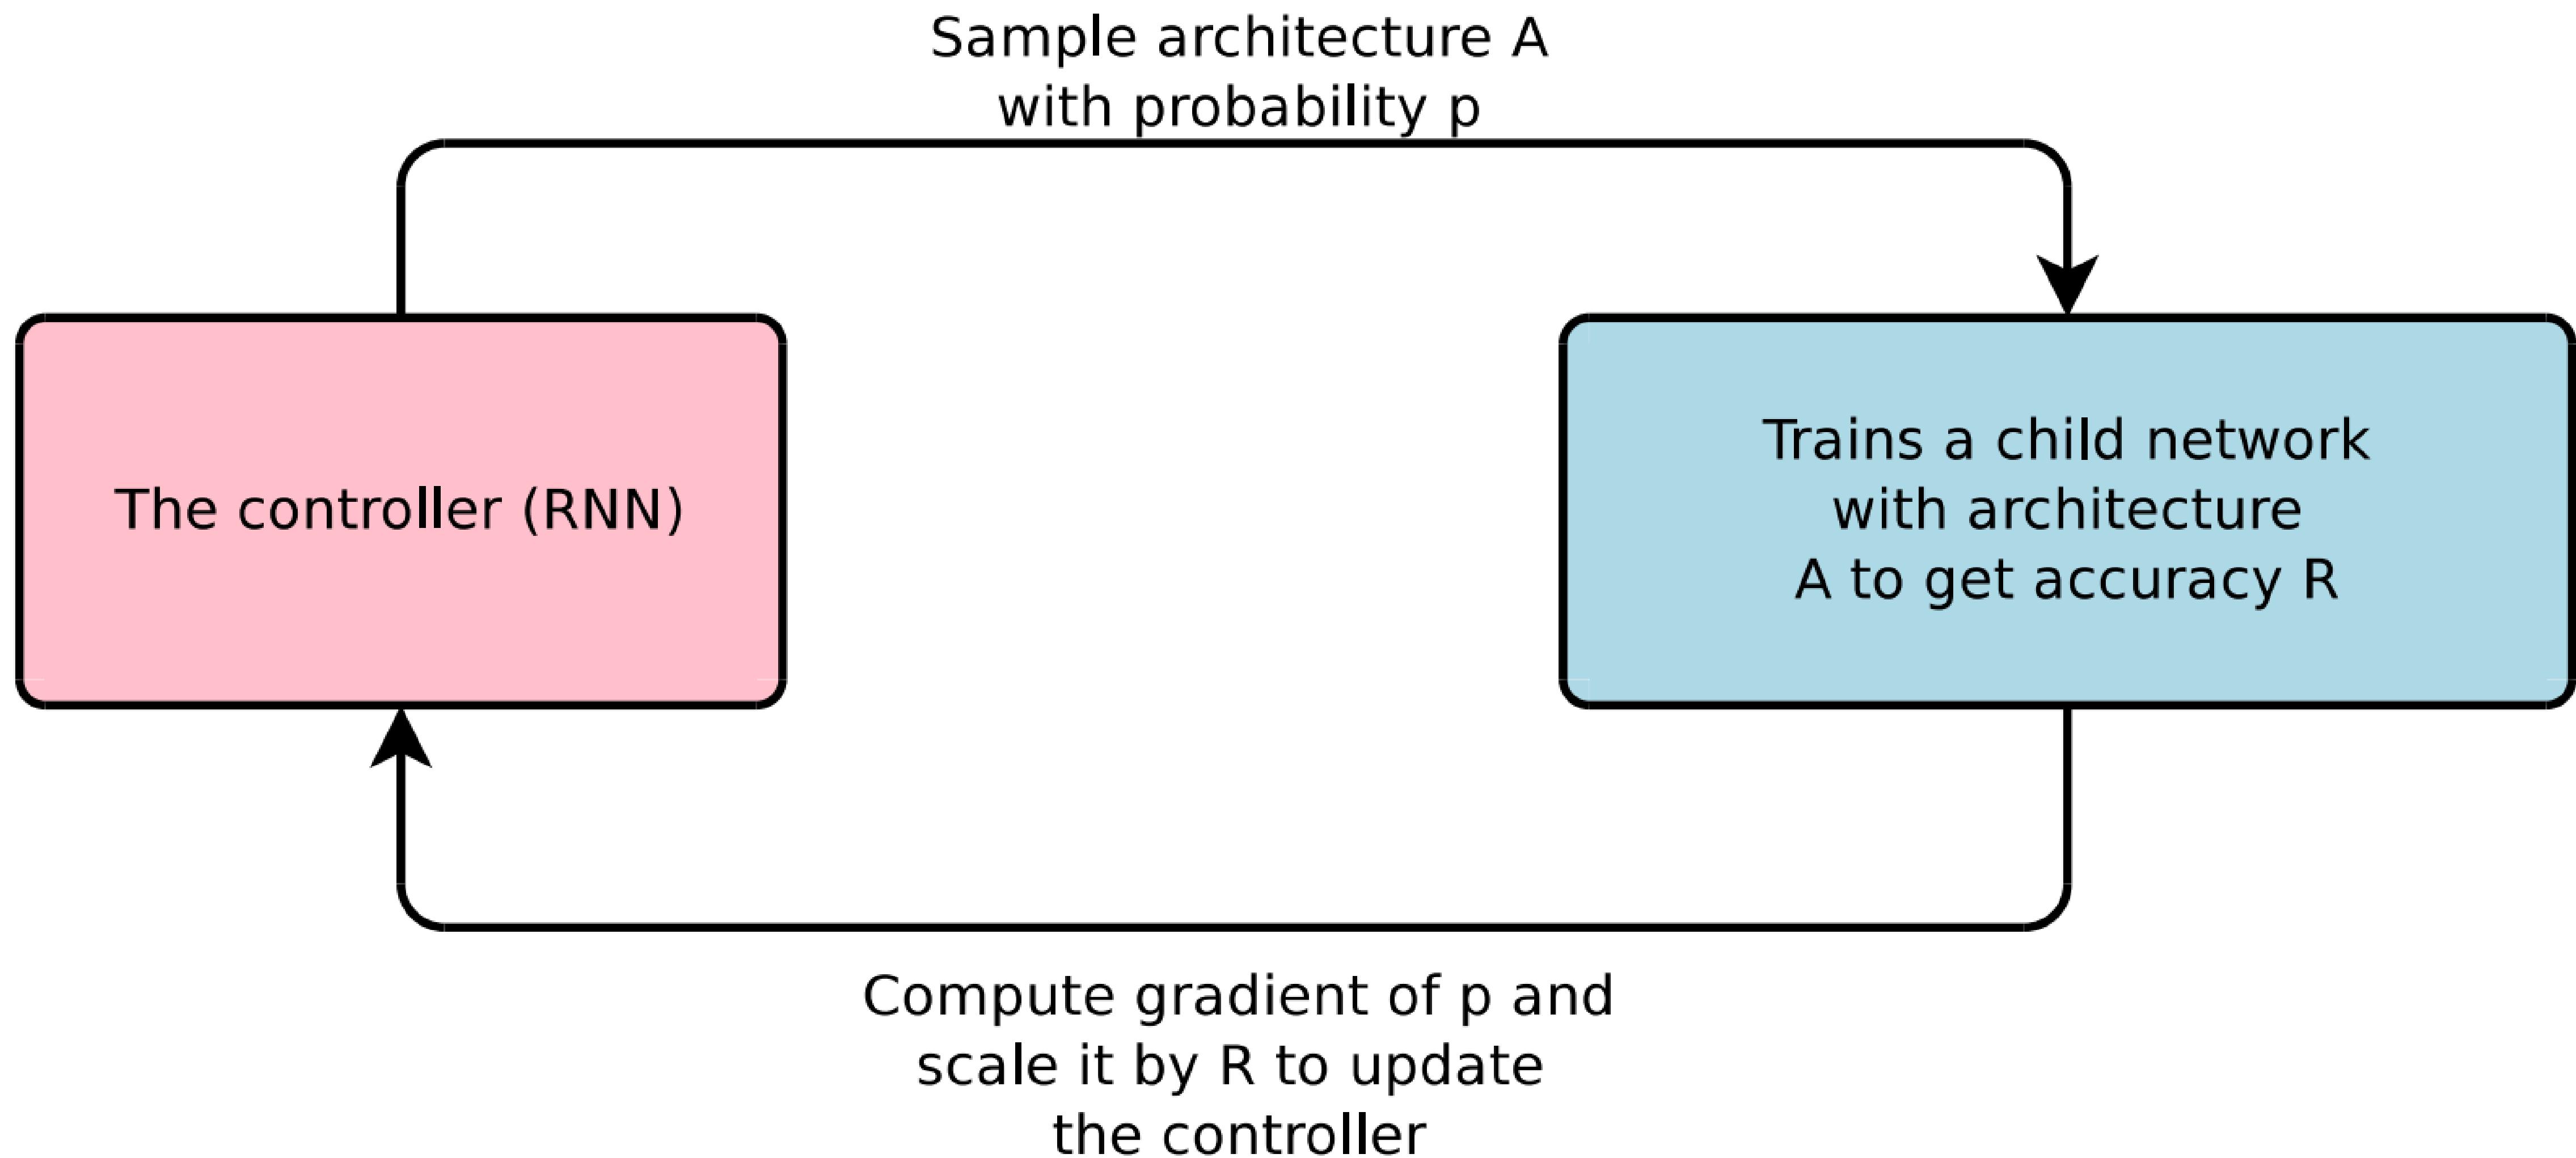


SENet (2017) | General idea implementation



Pre-attention

Neural Architecture Search | Overview



Neural Architecture Search | Example

